MAS DSE 260: Capstone Project

İlkay ALTINTAŞ, Ph.D.

Lecture 1: Setting up Your Data Process

Today's Topics

1. Reviewing where we are
2. STEP II: Designing the Data Acquisition and Preparation Pipelines
3. Report II Format : DUE 2/1/18

Process Roadmap (260 A)

- ✓ Step 1: Understanding the Challenge
  - ✓ REPORT 1: due 1/18
- Step 2: Designing the Data Acquisition and Preparation Pipelines
  - REPORT 2: due 2/1
- Step 3: Exploring Data
  - PRESENTATION 1: 2/3
  - REPORT 3: due 2/15
- Step 4: Defining Your Hypothesis and Minimum Viable Modeling Product
  - REPORT 4: due 3/1
- Step 5: Creating a Solution Architecture for Modeling and Optimization
  - PRESENTATION 2: 3/3
  - FINAL WINTER REPORT: due 3/16

Asking the Right Question

"A problem well defined is a problem half solved."

Charles F. Kettering

Define the Problem

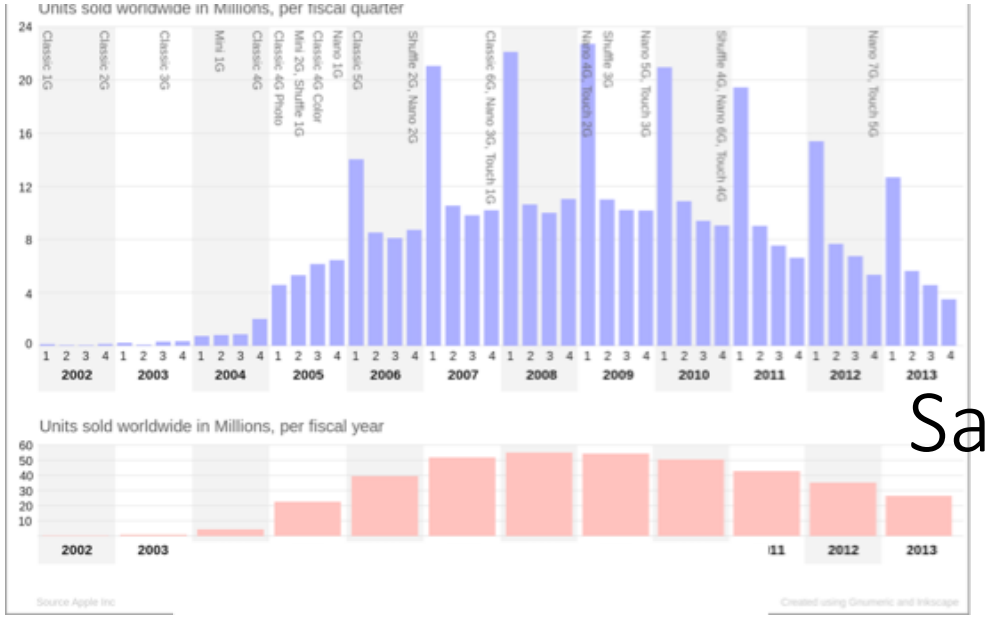

Sales figures

Call center logs

Evaluate a new product

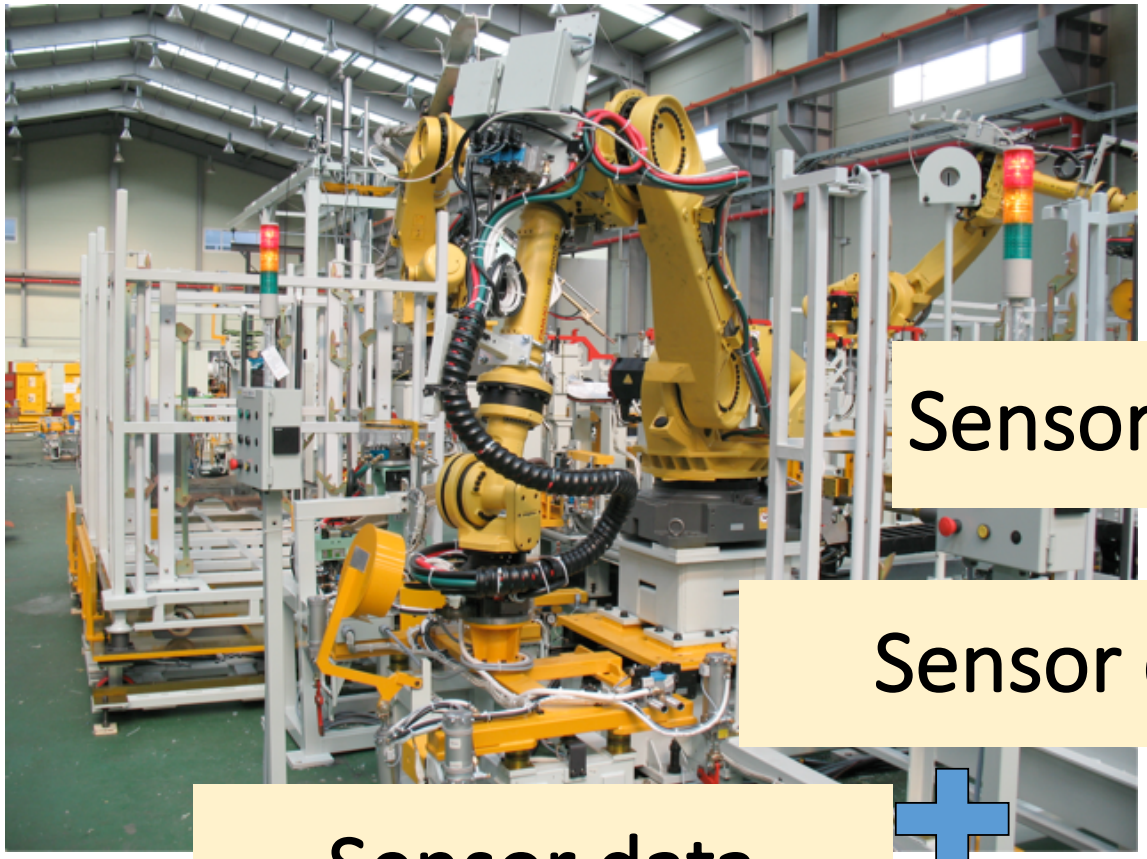

Detect equipment failure

Sensor data

Sensor data

Sensor data

Better targeted marketing

Customer data

+

Marketing data

1 Market Size and Growth

(ialtintas@ucsd.edu)

Assess the Situation

- Risks
- Benefits
- Contingencies
- Regulations
- Resources
- Requirements

Assess the Situation

Define Goals

Objectives

Criteria

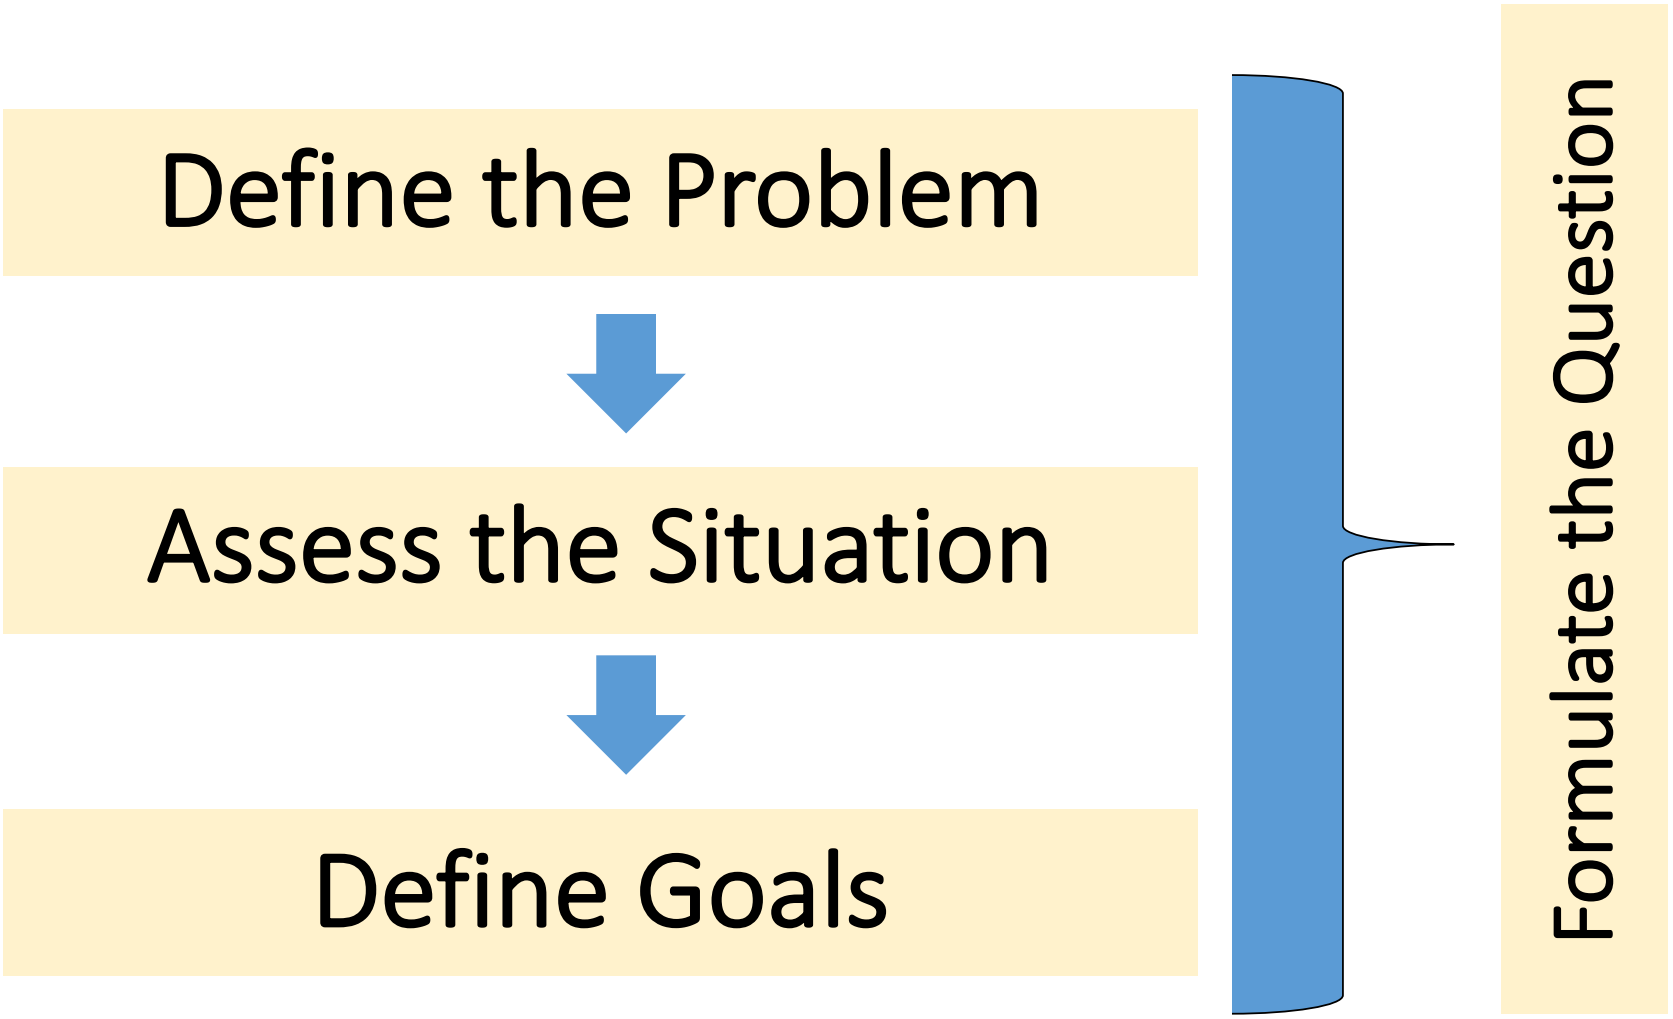
Define the Problem
Assess the Situation
Define Goals
Formulate the Question

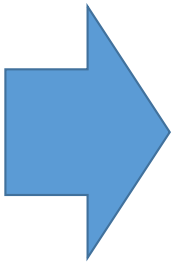

Data Science Process

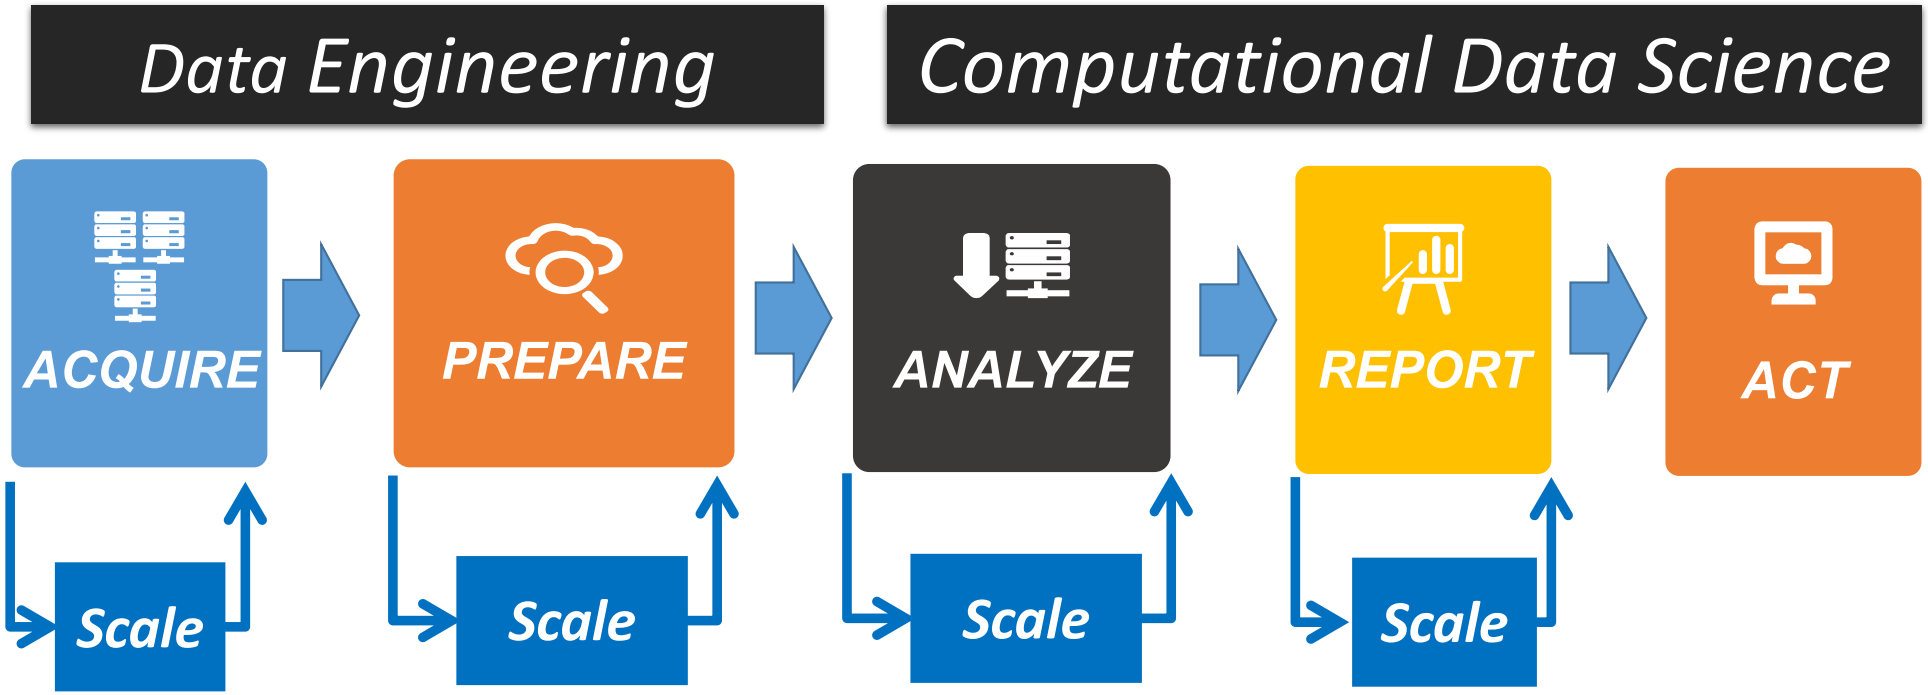

Many iterations and rollbacks between steps.

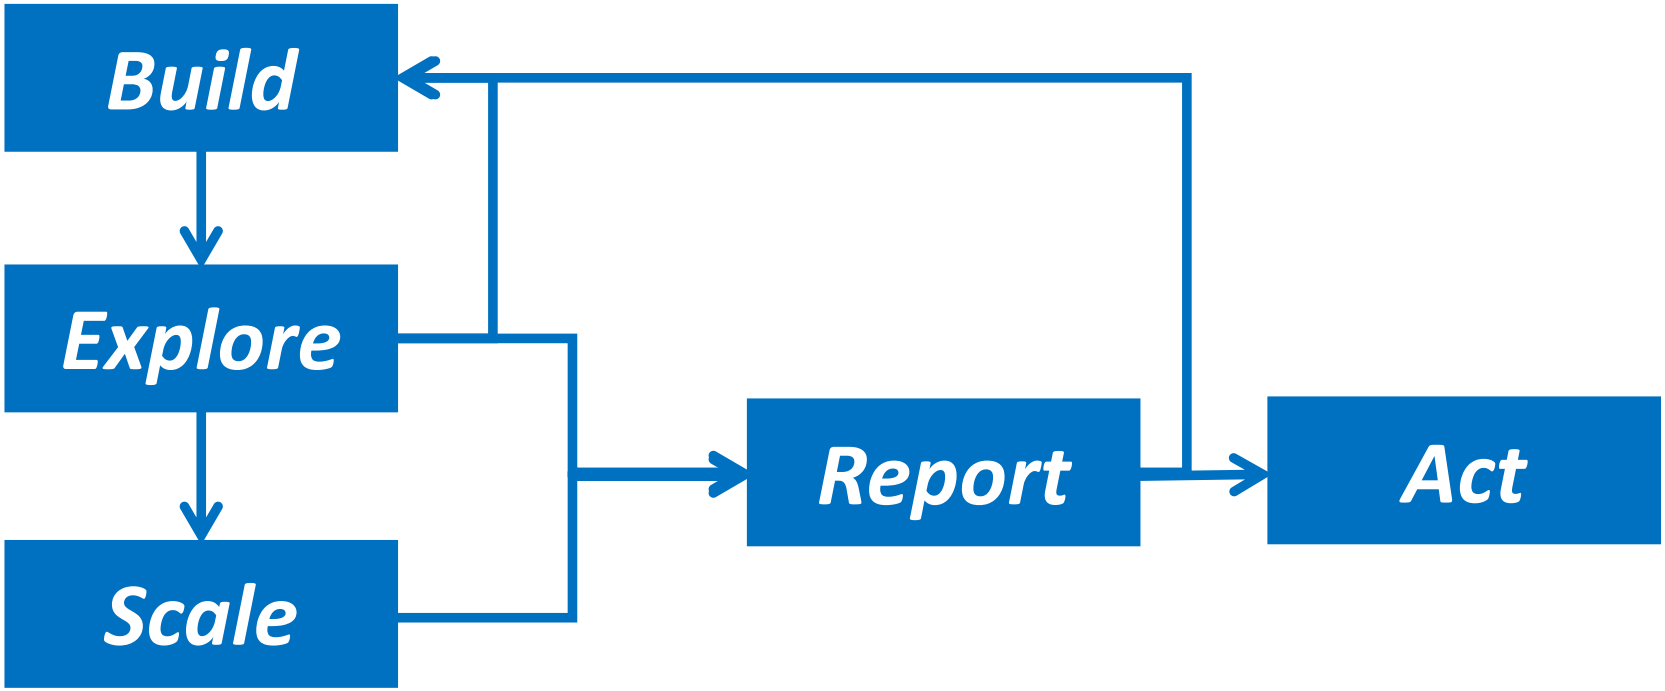
Build
Explore
Scale
Report
Act

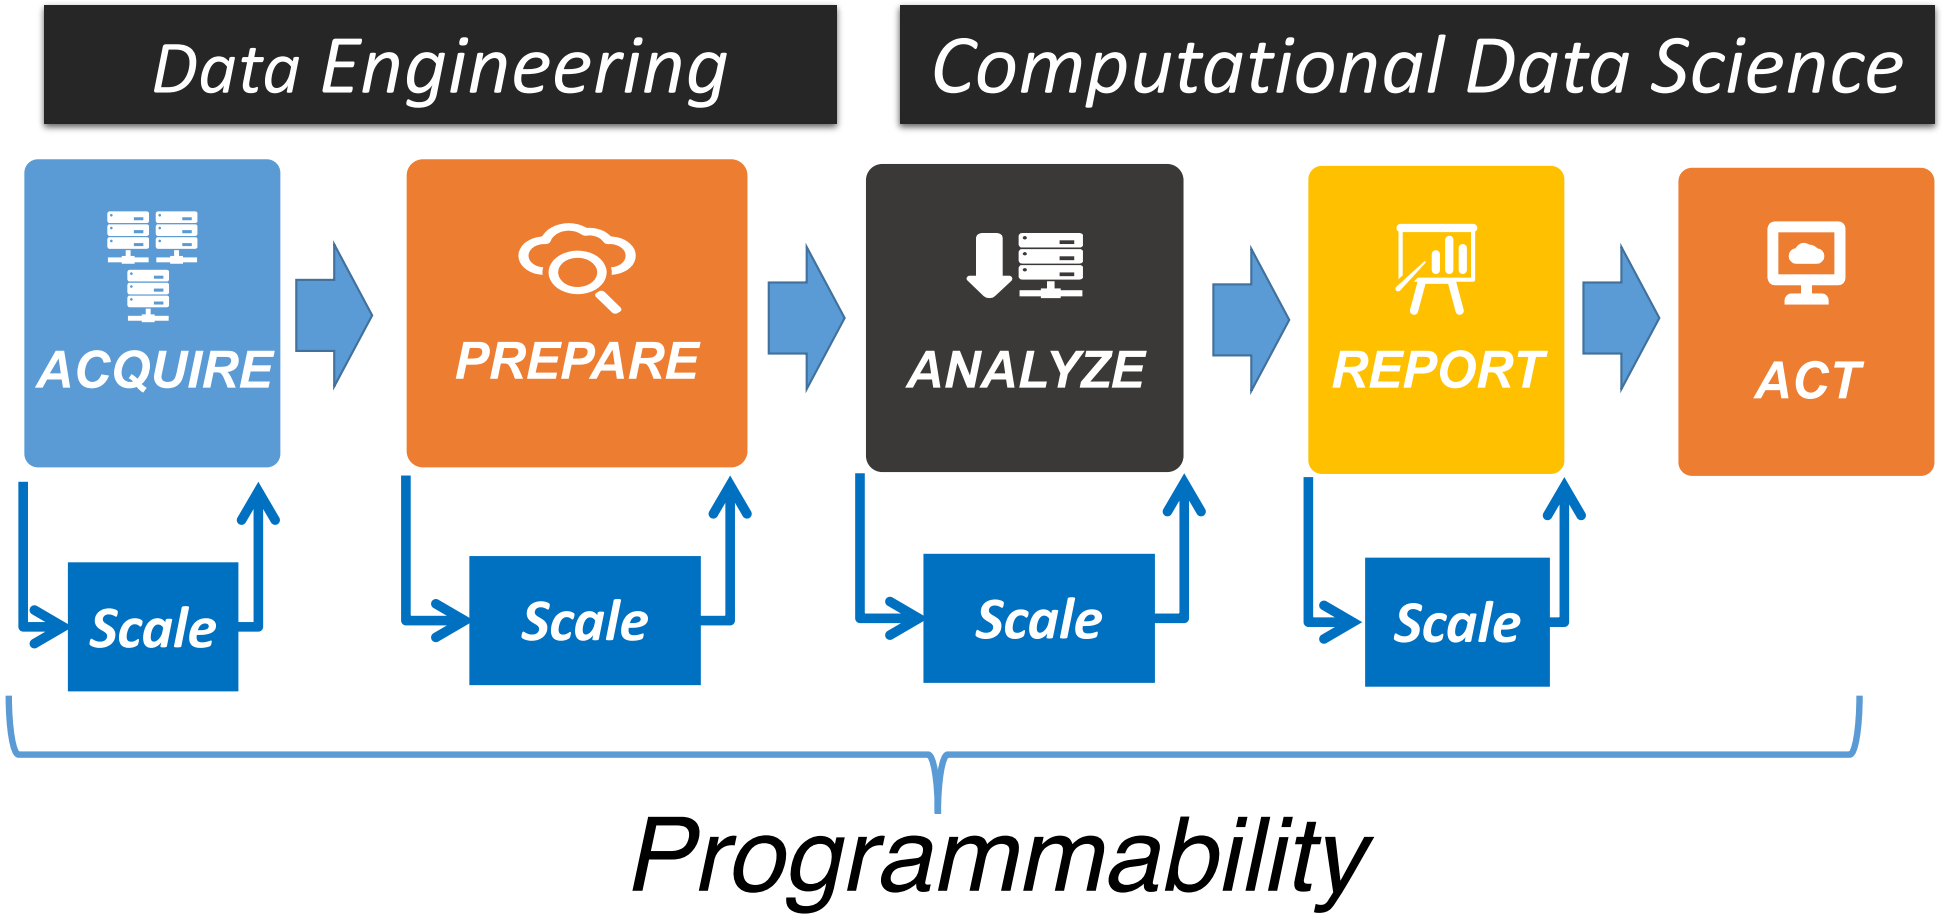
Data Engineering
Computational Data Science
ACQUIRE
PREPARE
ANALYZE
REPORT
ACT
Scale
Scale
Scale
Scale
Programmability

NEXT:

STEP 2: Designing the Data Acquisition and Preparation Pipelines

Why do we need this step?

- Data ingest and engineering is the basis for effective modeling
- Many quality issues with data that needs to be cleaned up
  - Not just missing values, but also deeper problems on uncertainty and bias related to data collection
- Defined data pipelines (potentially automated) help with data refresh and updates

If your data pipeline is strong, your modeling will be strong!

Objectives

- Understanding what data is available and useful to you
- Plan for data transport and locate your datasets in the right database, storage or analytics environment
- Model the data
- Identify data integration needs if applicable
- Set up a data pipeline that streamline (and potentially automates) data ingest and refresh
- Design programmatic access to data

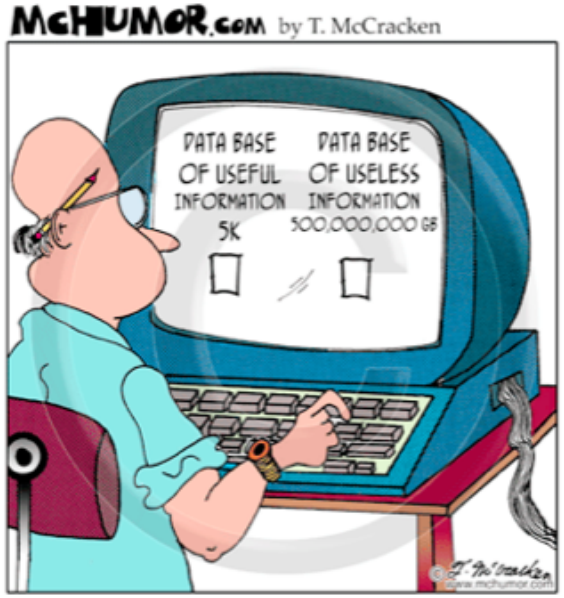

Planning for a Good Data Pipeline

- Focus on
  - where your datasets come from (data sources)
  - adding context and model data as you collect it
  - privacy and policy issues with the data
  - data quality assurance
  - where each dataset gets stored
  - how can pieces of data be combined
  - how each datasets get used
  - programmatic access to data

Using the PPODS Approach

- Each step in your data pipelines is a separate pod
- Define success metrics for calling a step done

Metrics for accountability should be built into the process.

Cost **Timeline**

Planning of deliverables

Purpose **Expectations**

Pod ➔ sub-process

Defined by:

- Purpose and goal
- Stakeholders
- Expectations
  - Key questions to be answered, production/consumption relationships, needs, dependencies, limits, …
- Contracts
  - Performance, economic, accuracy, policy, privacy, reproducibility, political, …
- Knowns
- Known unknowns

Step II Report Guidelines

- Title, team members and advisor(s)
- Sections:
  - Raw Data Sources
    - Summaries of each dataset description from report 1
    - Table for Dataset name, source location, destination in your data pipeline, data movement and processing scripts and notebooks, and data size
  - Data Exploration, Cleaning, Wrangling and Engineering
    - Data Exploration Summary
    - Data Preprocessing Approach
    - Approach for storing processed and/or integrated data
      - Processed dataset description for each processed dataset including why you want to process it that way
      - Table for processed data sets including processed data set name, input datasets, link to the processing scripts and notebooks, and provisional data size
    - Approach for Feature Engineering and Data Modeling
      - Summary of feature sets
      - Table for feature set including links to input datasets, feature engineering scripts and notebooks, and provisional data size
    - Approach for Data Access
      - Initial design for data querying interfaces
      - Justification for manual vs. programmatic access
  - Data Pipeline
    - Description of the needs, approach, and data access and refresh frequency
    - Logical diagram showing major data pipeline components for data sources and sinks
  - Set up for your data environment
    - Cloud vs. local, database vs. flat files, etc.
  - Bullets for each team member's individual contributions in Step 2
  - Any major updates to Step 1 as a result of data pipeline step
- Keep it to 5-8 pages
- Due date: 2/1/2018 midnight

You will have more time to iterate. We are doing the first draft. Next, are some reminders on what to start focusing on.

Focus 1: Acquiring Data

After you are done with this part, you should be able to..

- List techniques and technologies to access and retrieve the data you need
- Describe an example scenario related to your problem accessing your datasets

Data can come from a variety of sources using different technologies

Data comes from many places…

…with many ways to access it.

Traditional databases

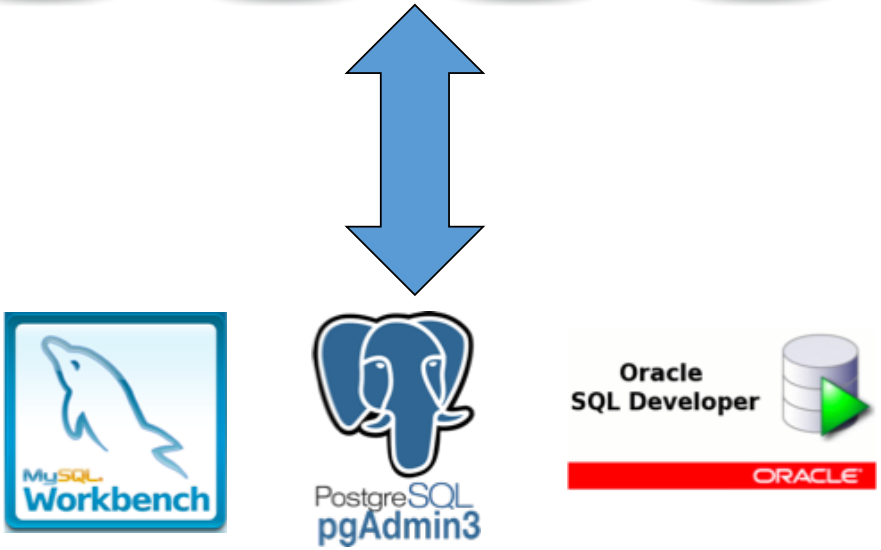

SQL and query browsers

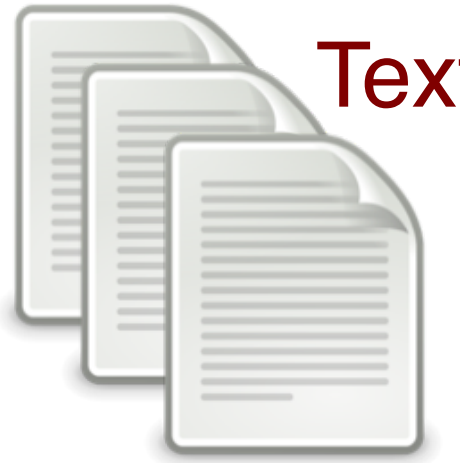

Text files

Scripting languages

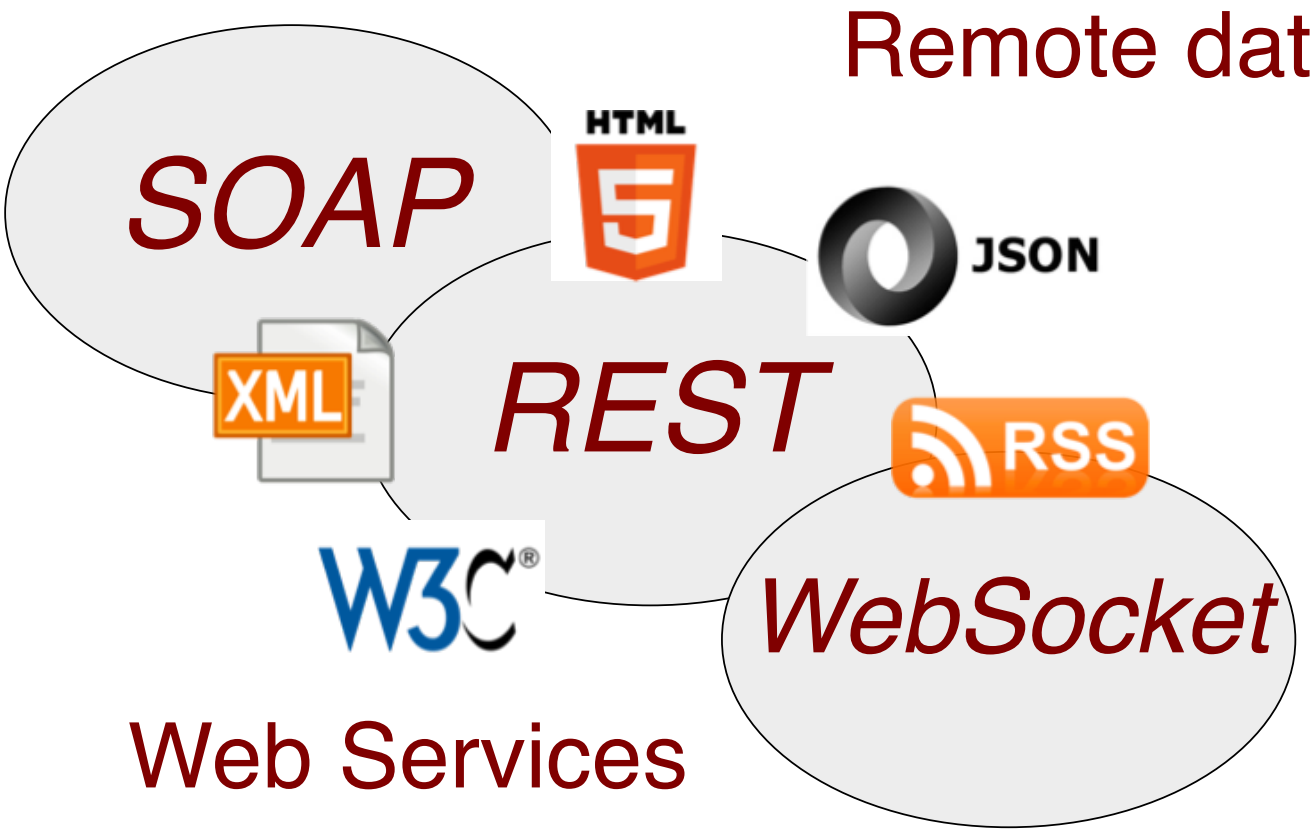
Remote data
SOAP
HTML
JSON
XML
REST
RSS
W3C
WebSocket
Web Services

NoSQL storage

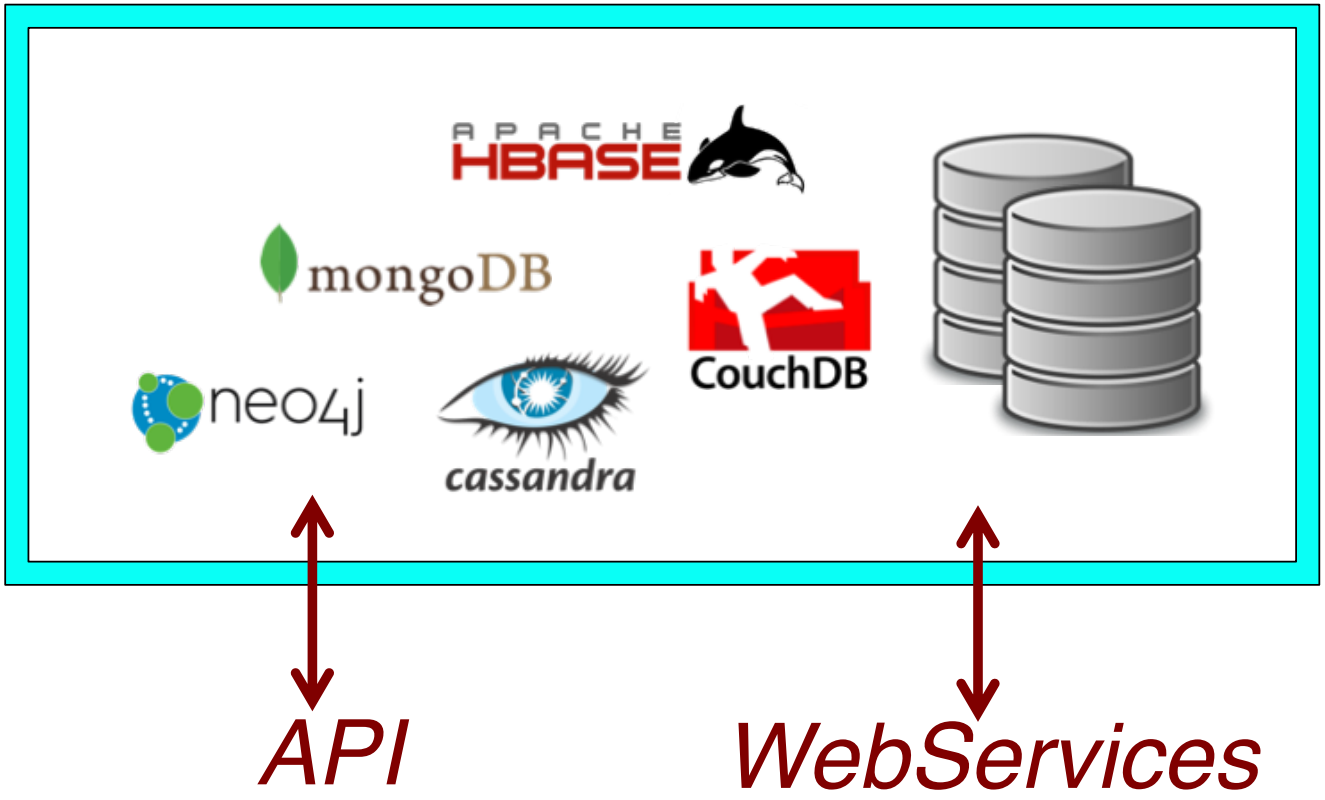

Acquiring Data From WIFIRE

Historical weather

SQL 

Current weather

WebSocket 

Real-time tweets near fires

REST 

Traditional databases

SQL and query browsers

Remote data

Web Services

Text files

Scripting languages

NoSQL storage

Web Services

Programming Interfaces

Focus 2: Exploring Data

Why Explore?

Goal: Understand your data

Why Explore?

Outliers

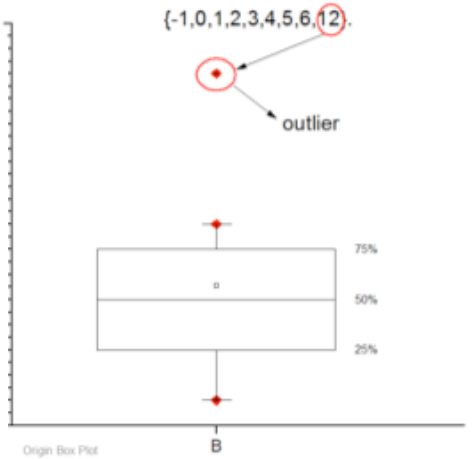

General trends

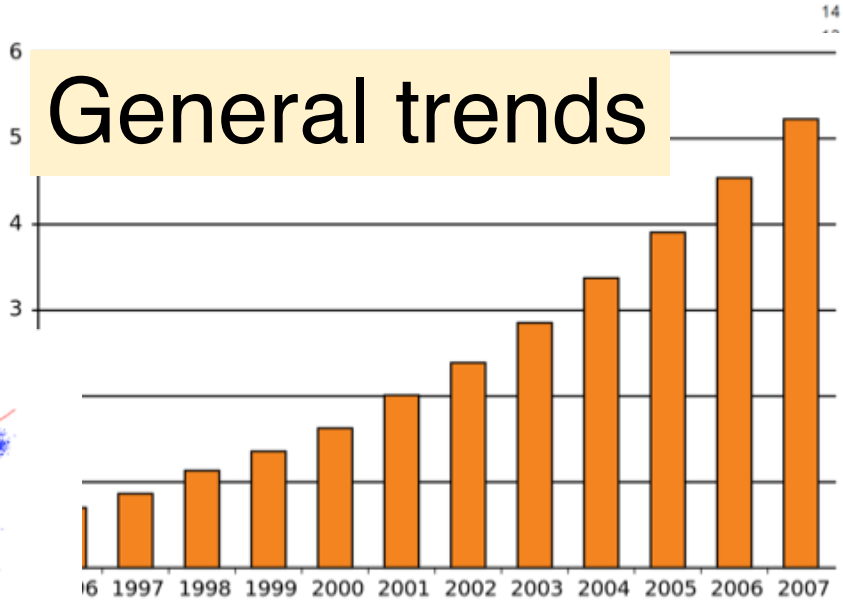

Correlations

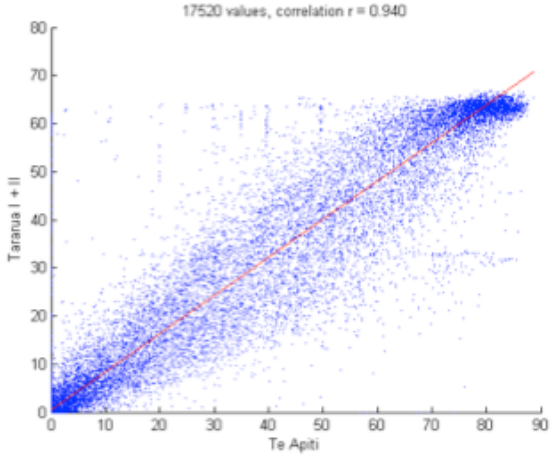

Describe Your Data

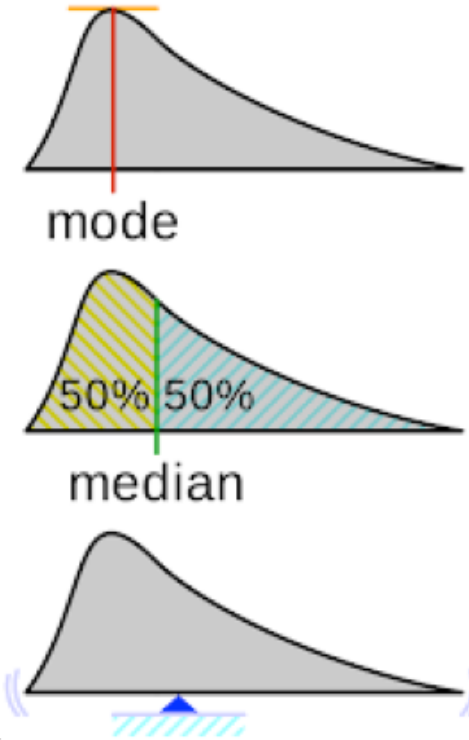

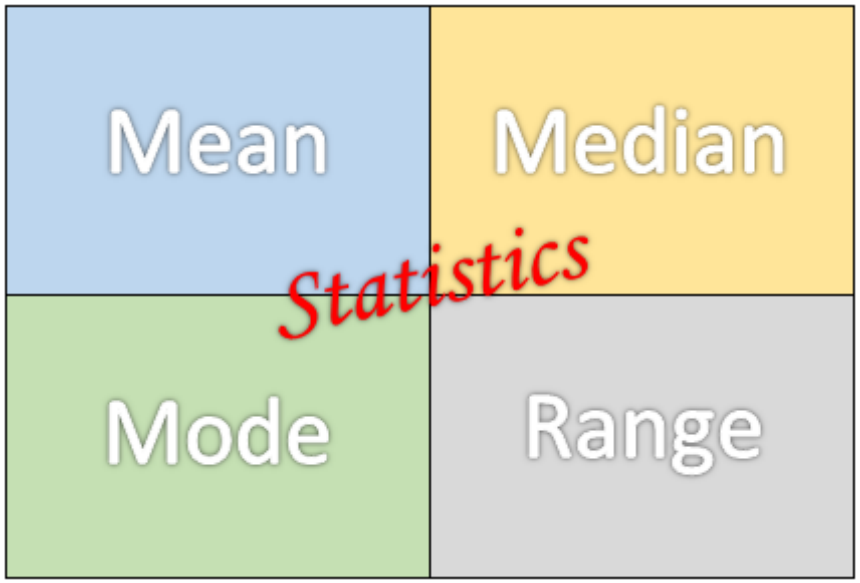

Visualize Your Data

Heat maps

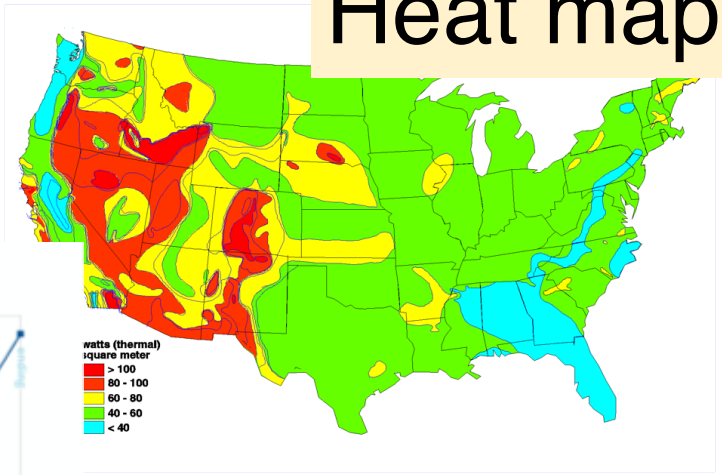

Histogram

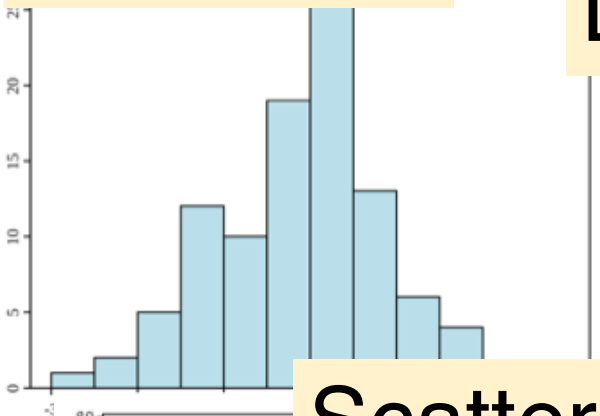

Line graphs

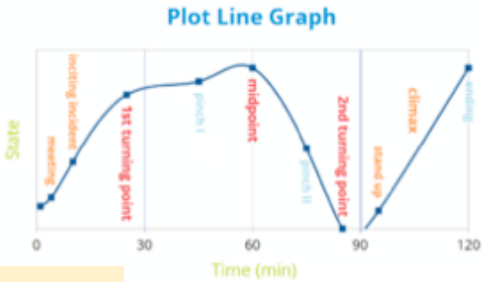

Scatter plots

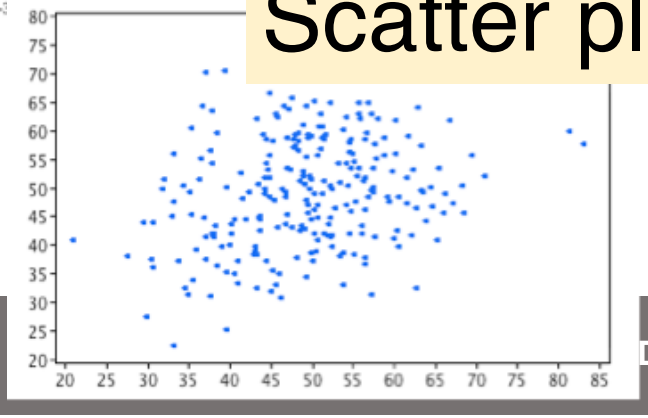

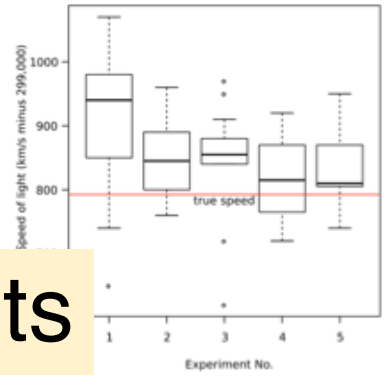

Boxplots

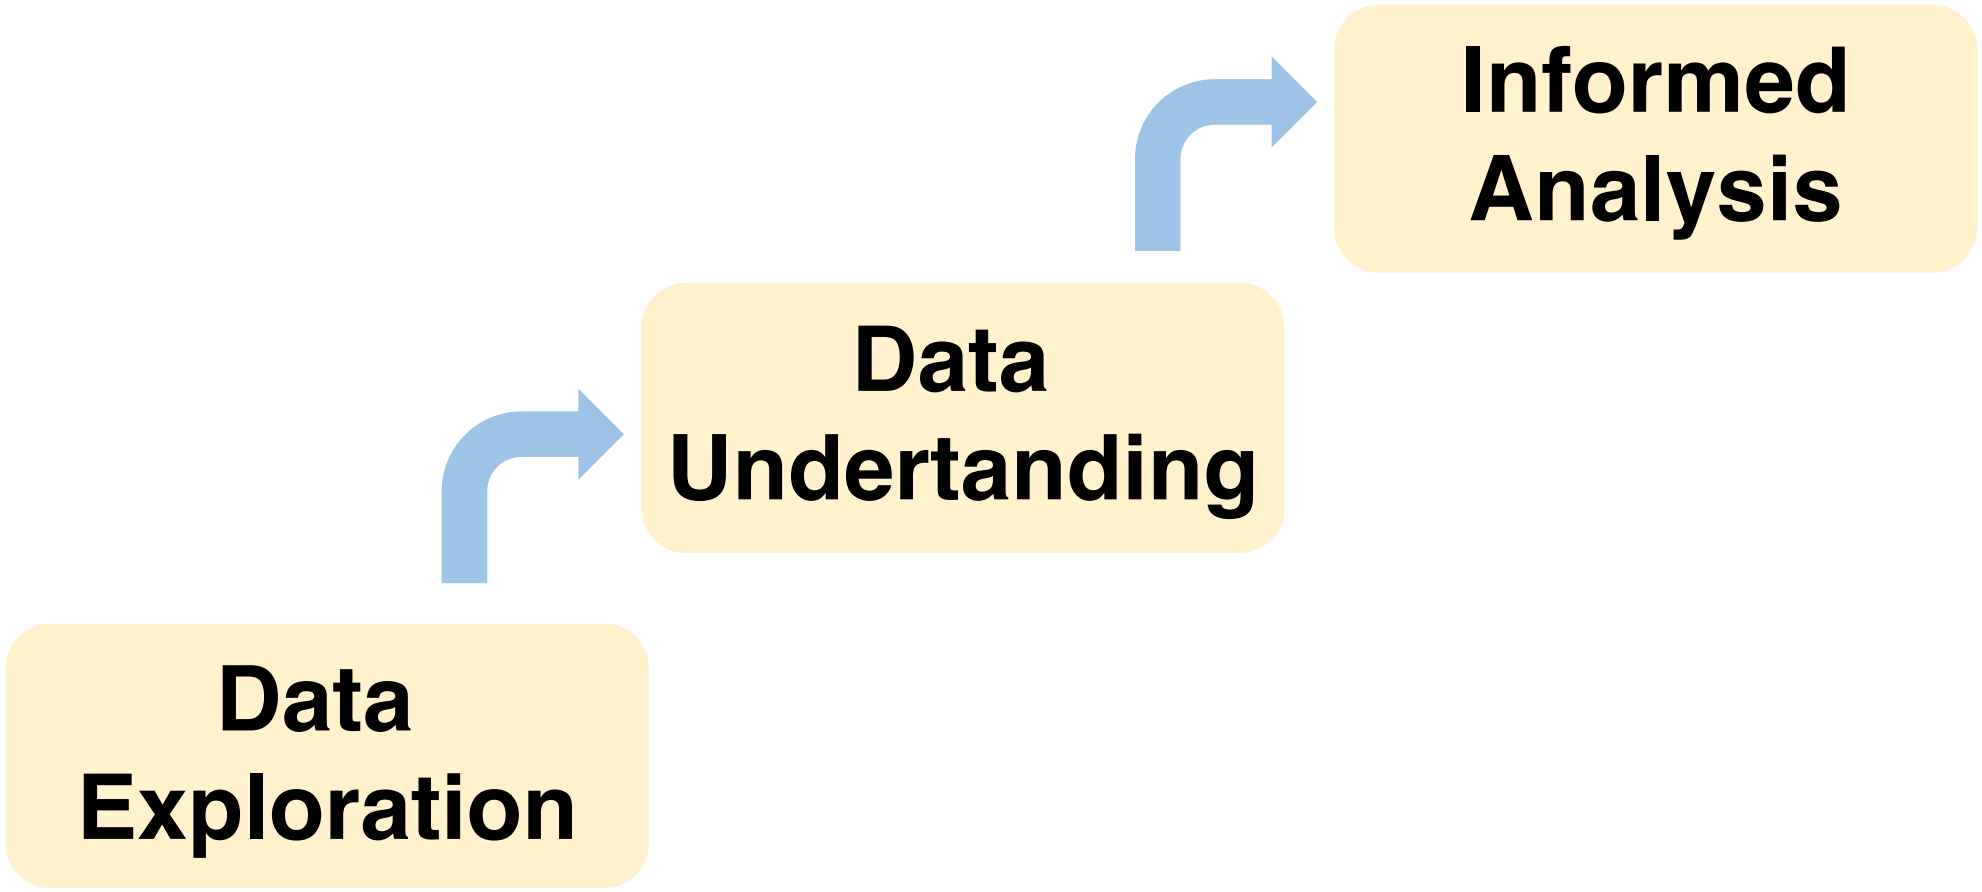
Informed Analysis
Data Undertanding
Data Exploration

Focus 3: Pre-processing Data

Describe what is needed to transform raw data to data that can be used for analysis.

Clean  Transform

Real-world data is messy!

- Inconsistent values
- Duplicate records
- Missing values
- Invalid data
- Outliers

Data Quality Issues

Addressing Data Quality Issues

Remove data with missing values

Generate best estimate for invalid values

Merge duplicate records

Remove outliers

Addressing Data Quality Issues

Remove data with missing values

Generate best estimate for invalid values

Merge duplicate records

Remove outliers

Domain Knowledge

Getting Data in Shape

Data Munging

Data Preprocessing

Data Wrangling

Data Munging

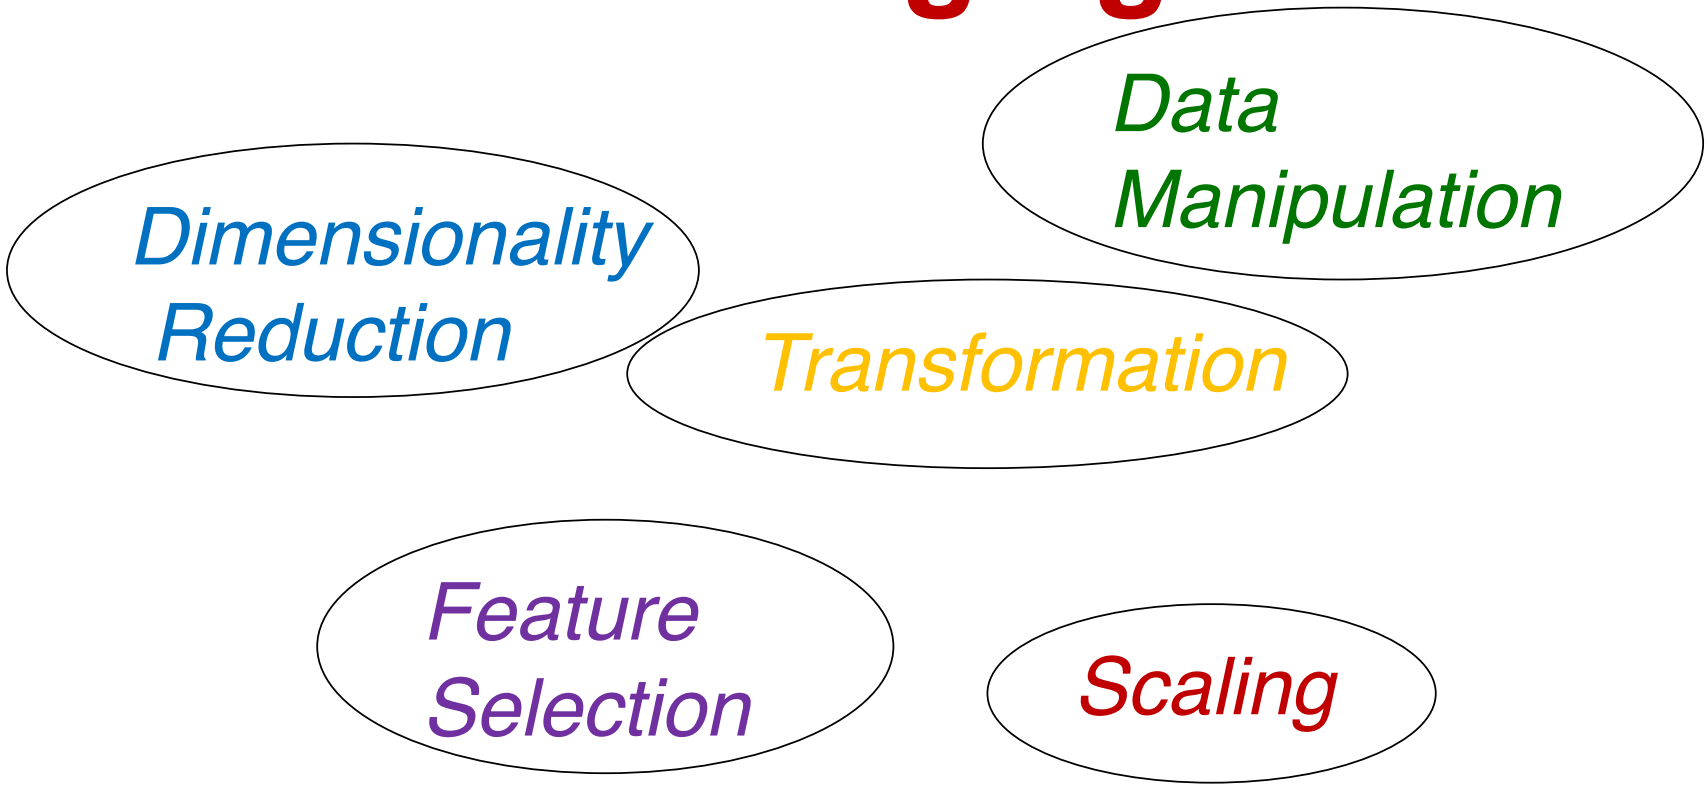

Scaling

Weight

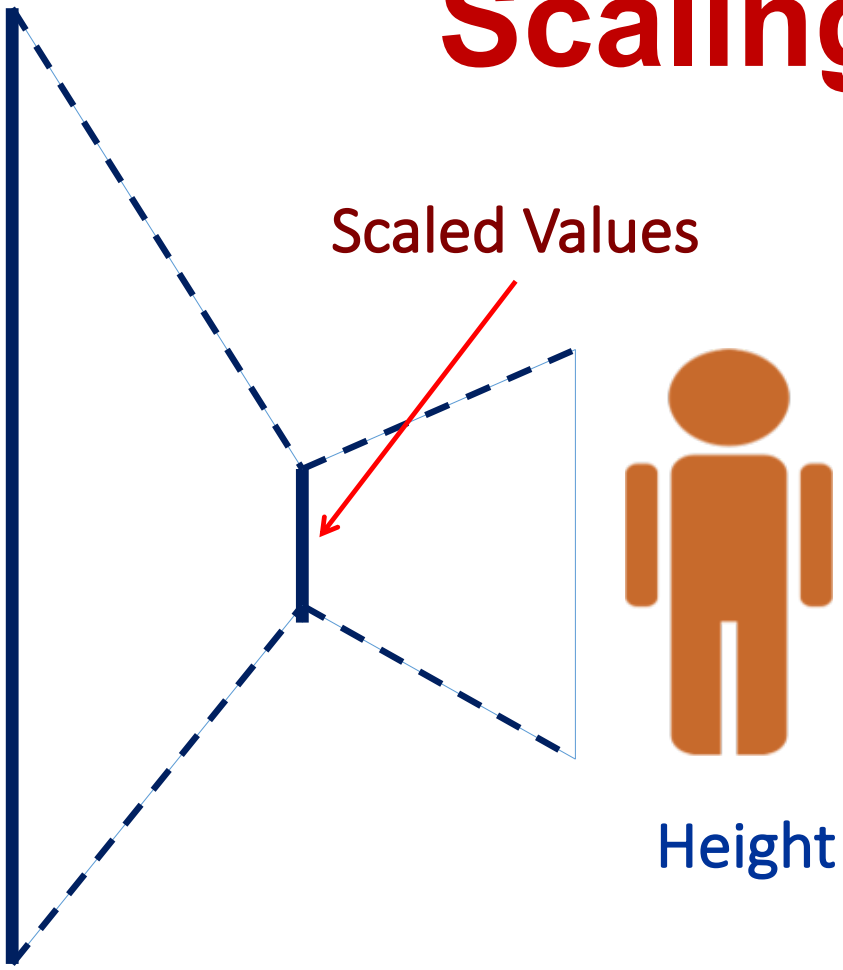

Height

Transformation

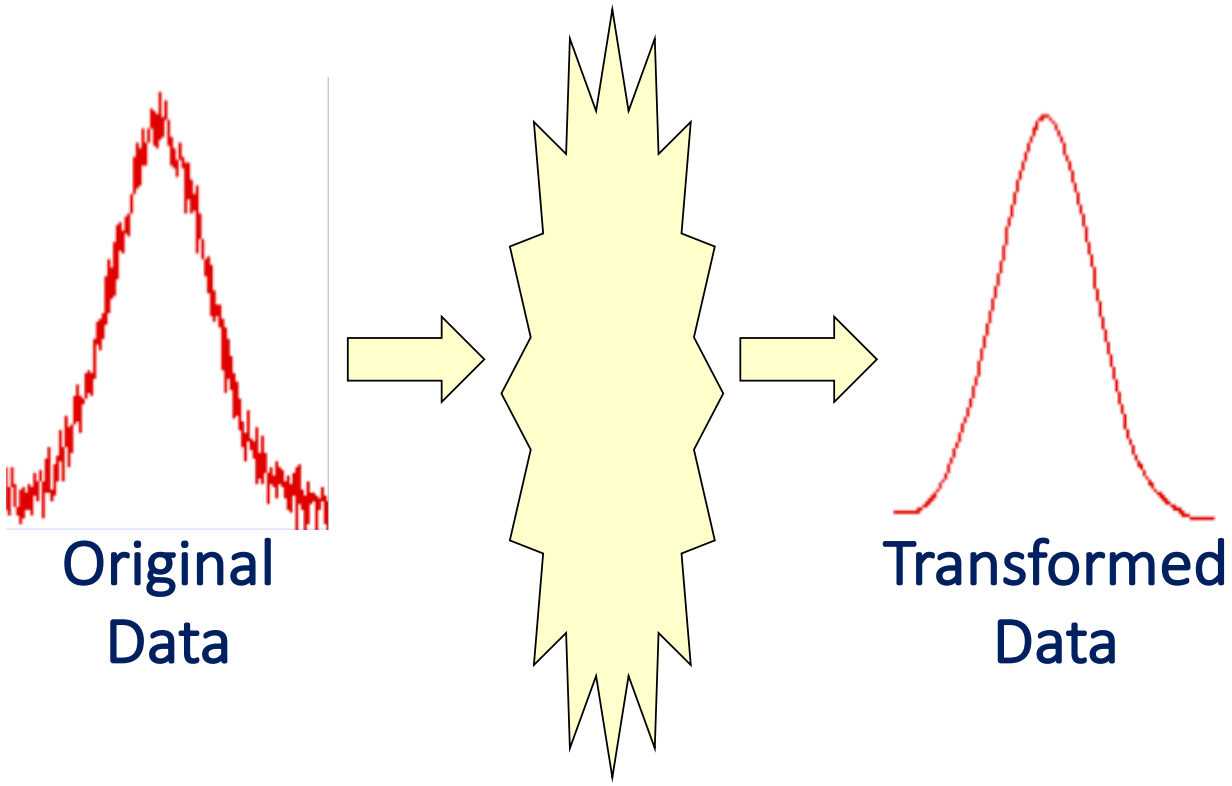

Feature Selection

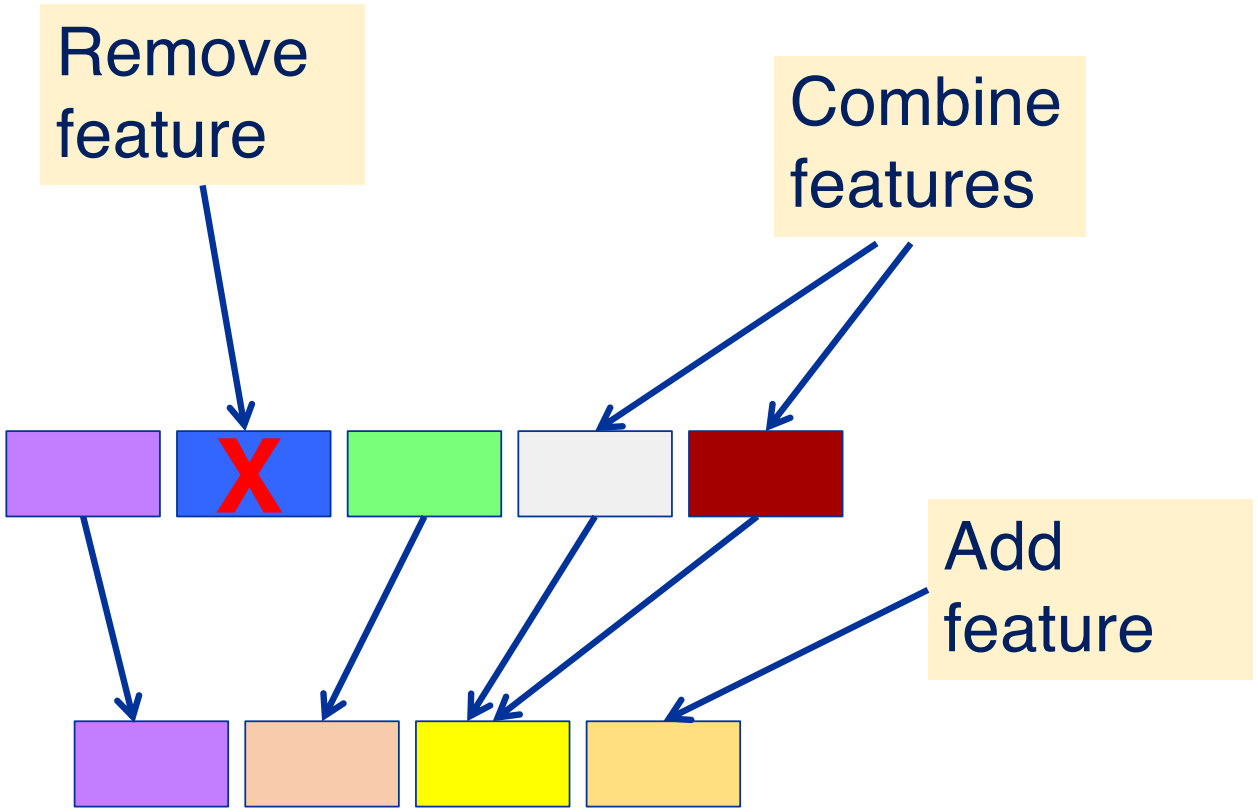

Dimensionality Reduction

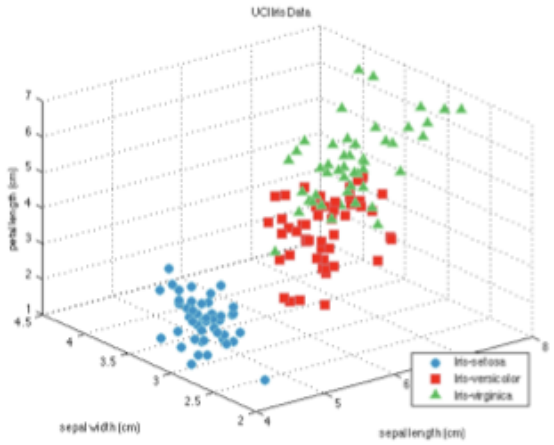

3D

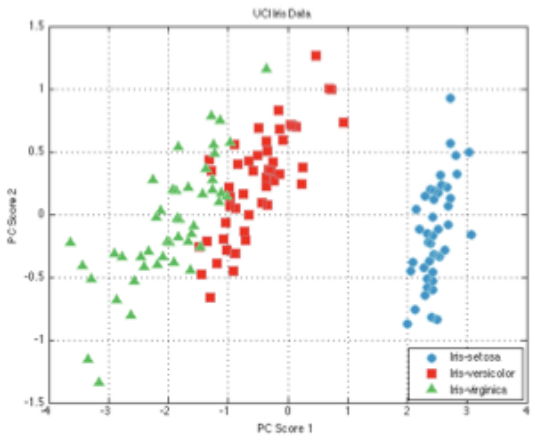

2D

Data Manipulation

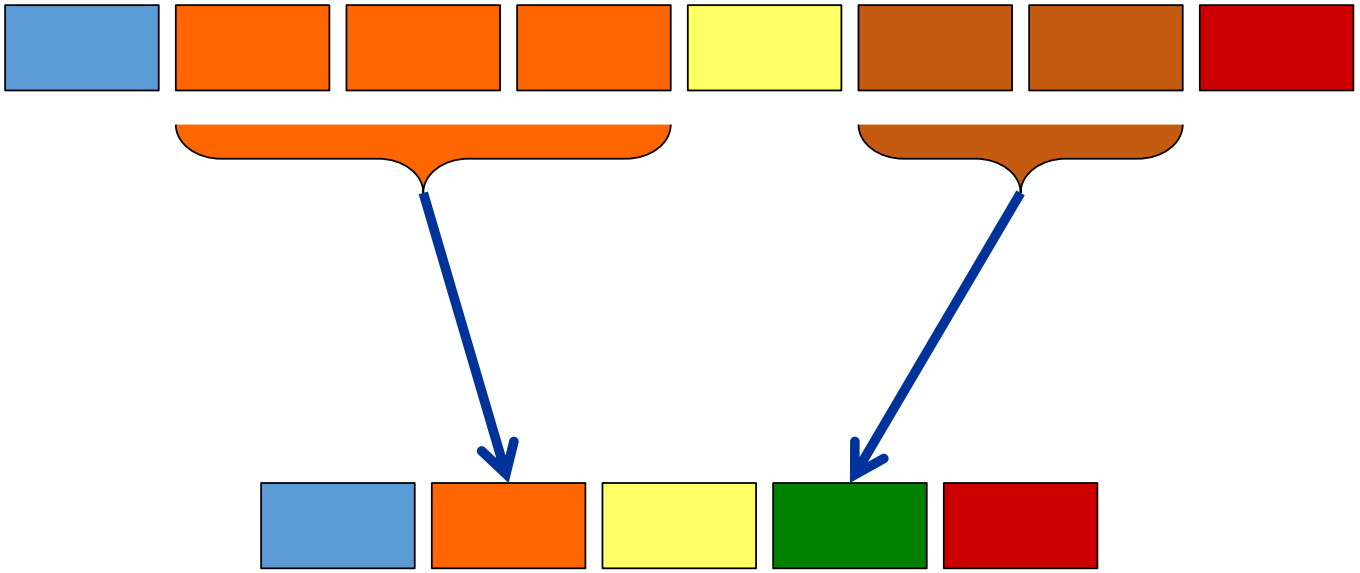

Always Remember!

Garbage in = Garbage out

Data preparation is very important for meaningful analysis!

Focus 4: Designing Your Data Pipelines

Draft solution architecture to turn data operations into an ongoing and accountable process.

Solution Architecture Diagram + Step by step flow diagram + Environment

Consider…

- Batch vs. streaming or both
- Local vs. cloud or both
- Database vs. object storage or both
- Manual vs. automated
- Refresh frequency
- Data provenance

A UNIX PIPE

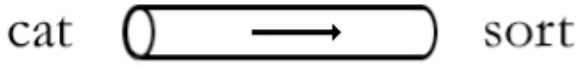

A UNIX pipe provides one-way communication between two processes on the same computer

Batch Example Data Pipeline

Represents a large number of applications.

Example MapReduce Application: WordCount

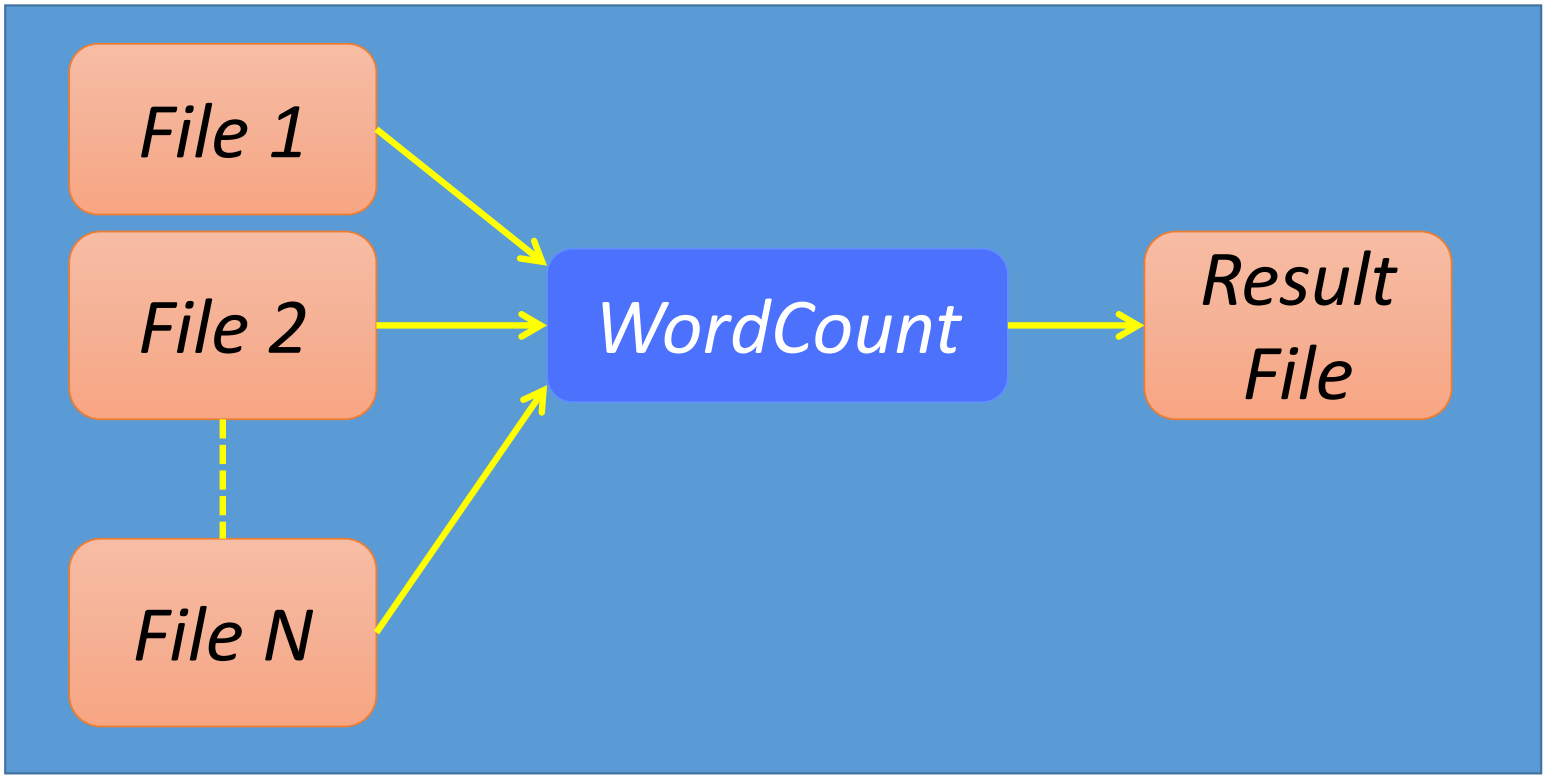

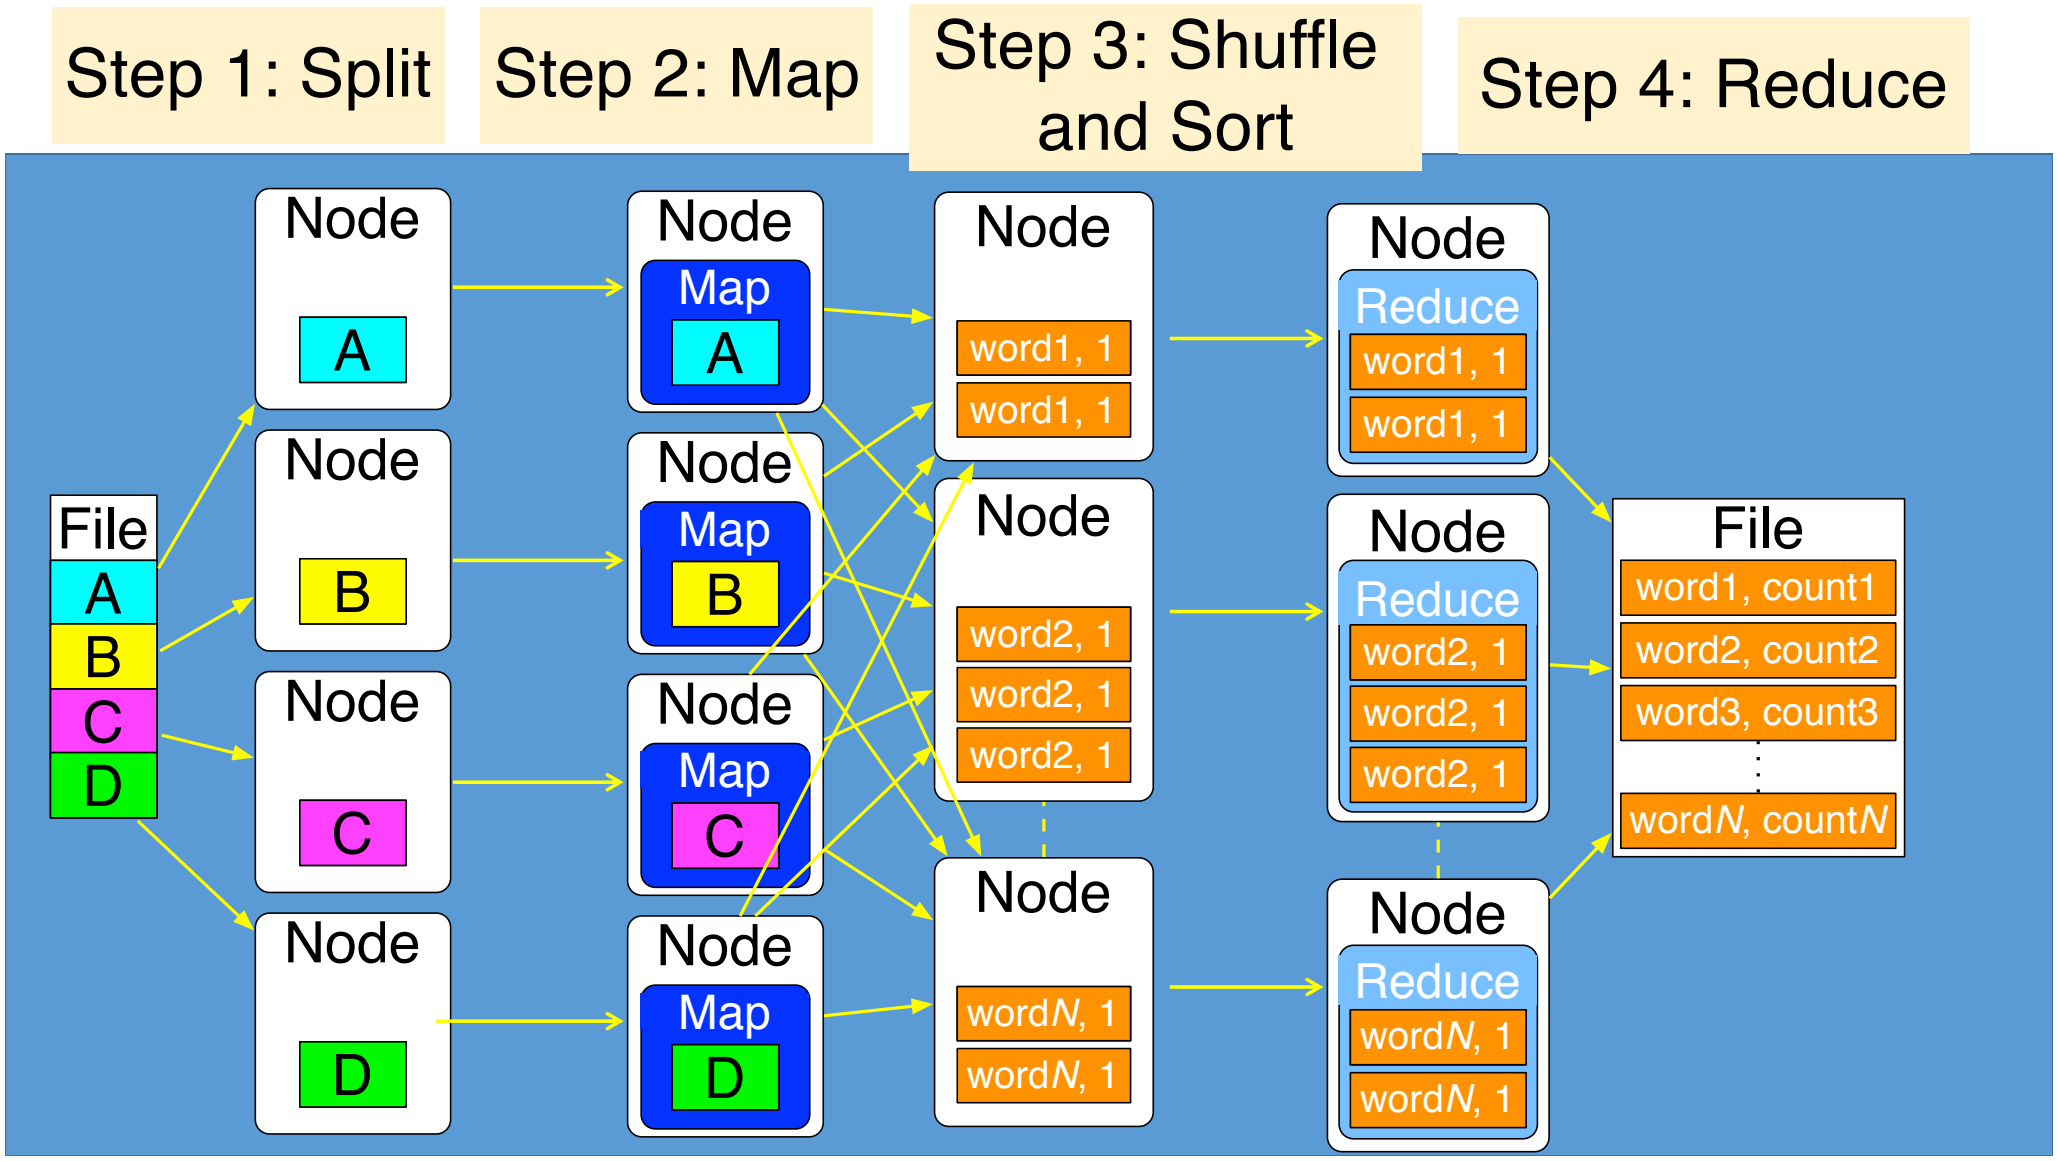
Step 1: Split
Step 2: Map
Step 3: Shuffle and Sort
Step 4: Reduce
File
A
B
C
D
Node
A
Node
B
Node
C
Node
D
Node
Map
A
Node
Map
B
Node
Map
C
Node
Map
D
Node
word1, 1
word1, 1
Node
word2, 1
word2, 1
word2, 1
Node
wordN, 1
wordN, 1
Node
Reduce
word1, 1
word1, 1
Node
Reduce
word2, 1
word2, 1
word2, 1
Node
Reduce
wordN, 1
wordN, 1
File
word1, count1
word2, count2
word3, count3
wordN, countN

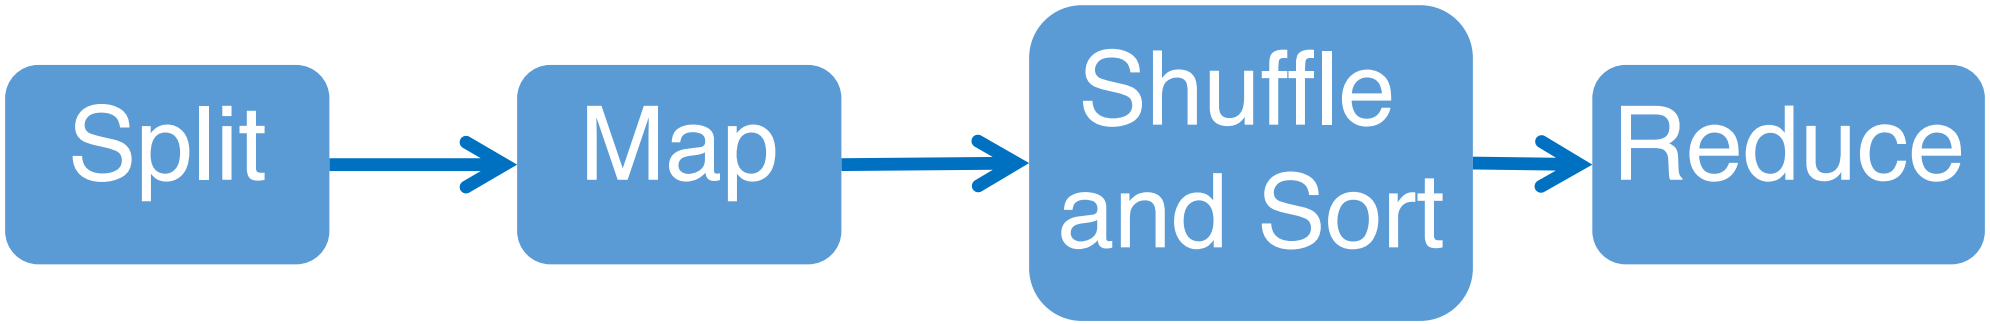
Split
Map
Shuffle and Sort
Reduce


Map
Shuffle and Sort
Reduce

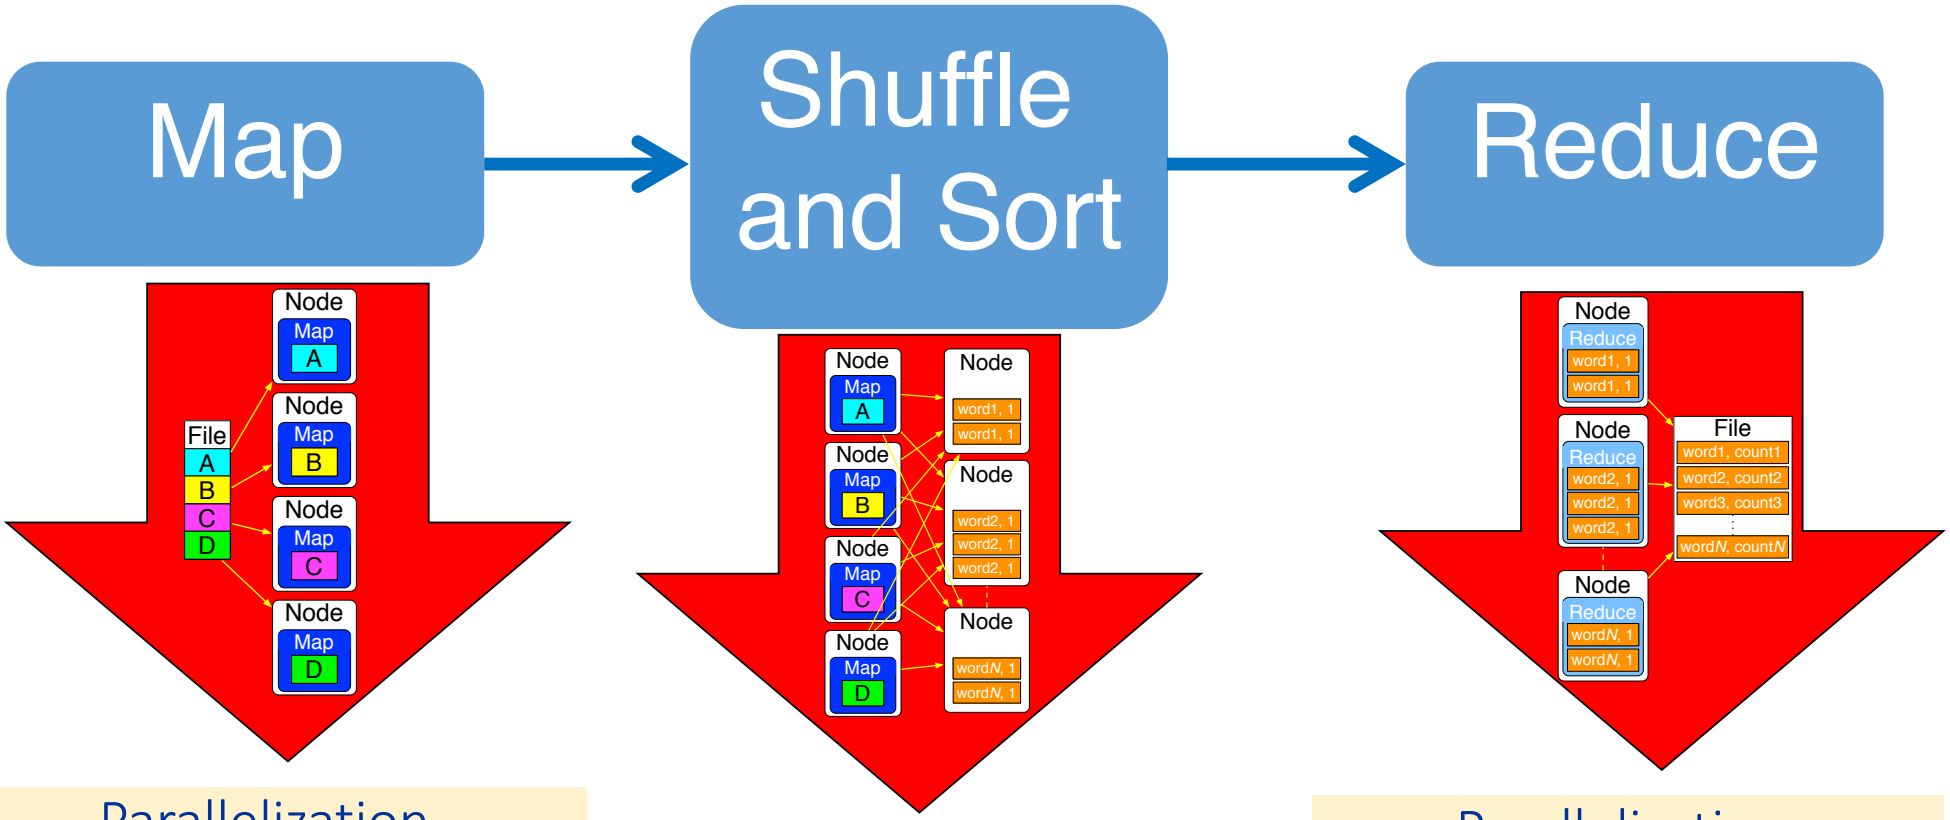

Parallelization over the input

Parallelization over intermediate data

Parallelization over data groups

Streaming Data

What is a Data Stream?

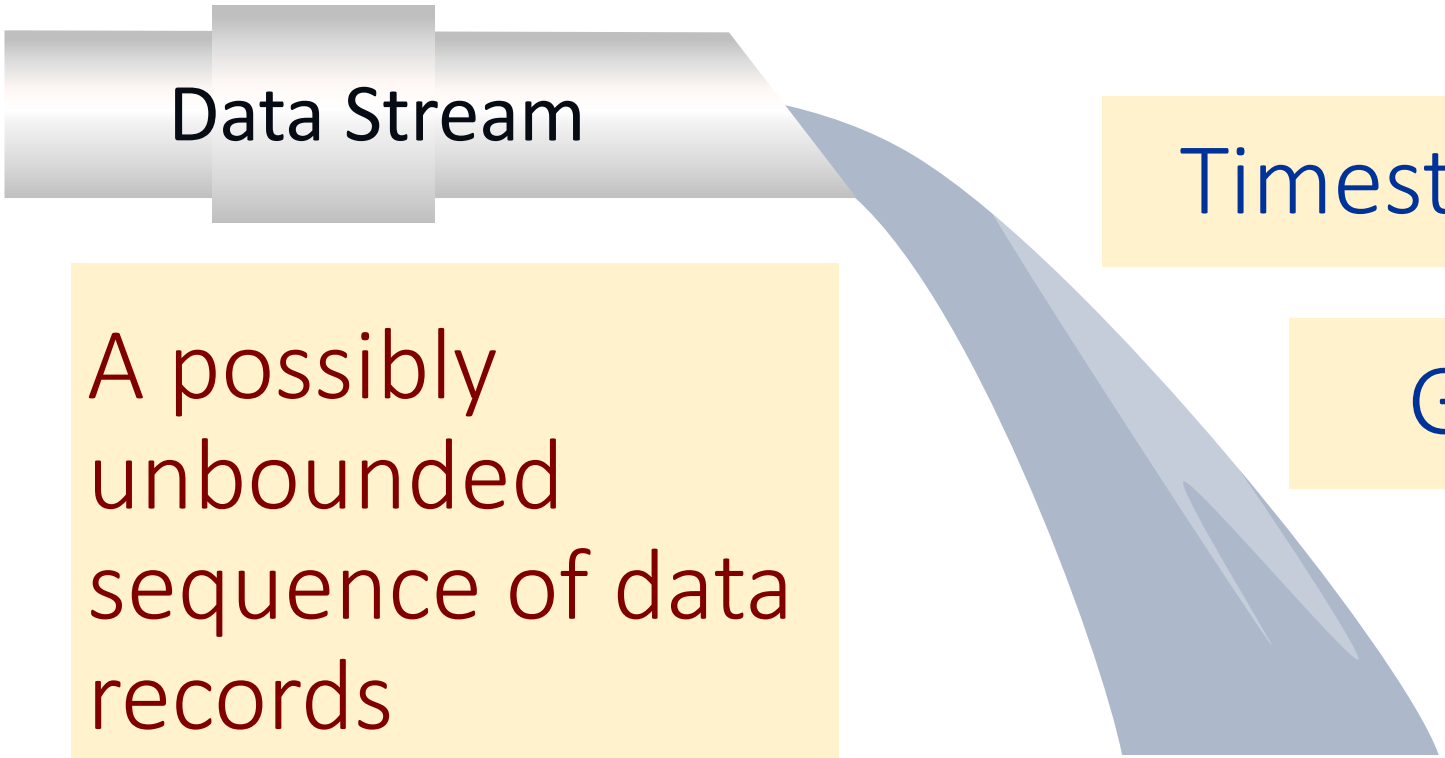

Timestamped

Geo-tagged

A possibly unbounded sequence of data records

Synchronized sequence of events

Streaming Data Systems

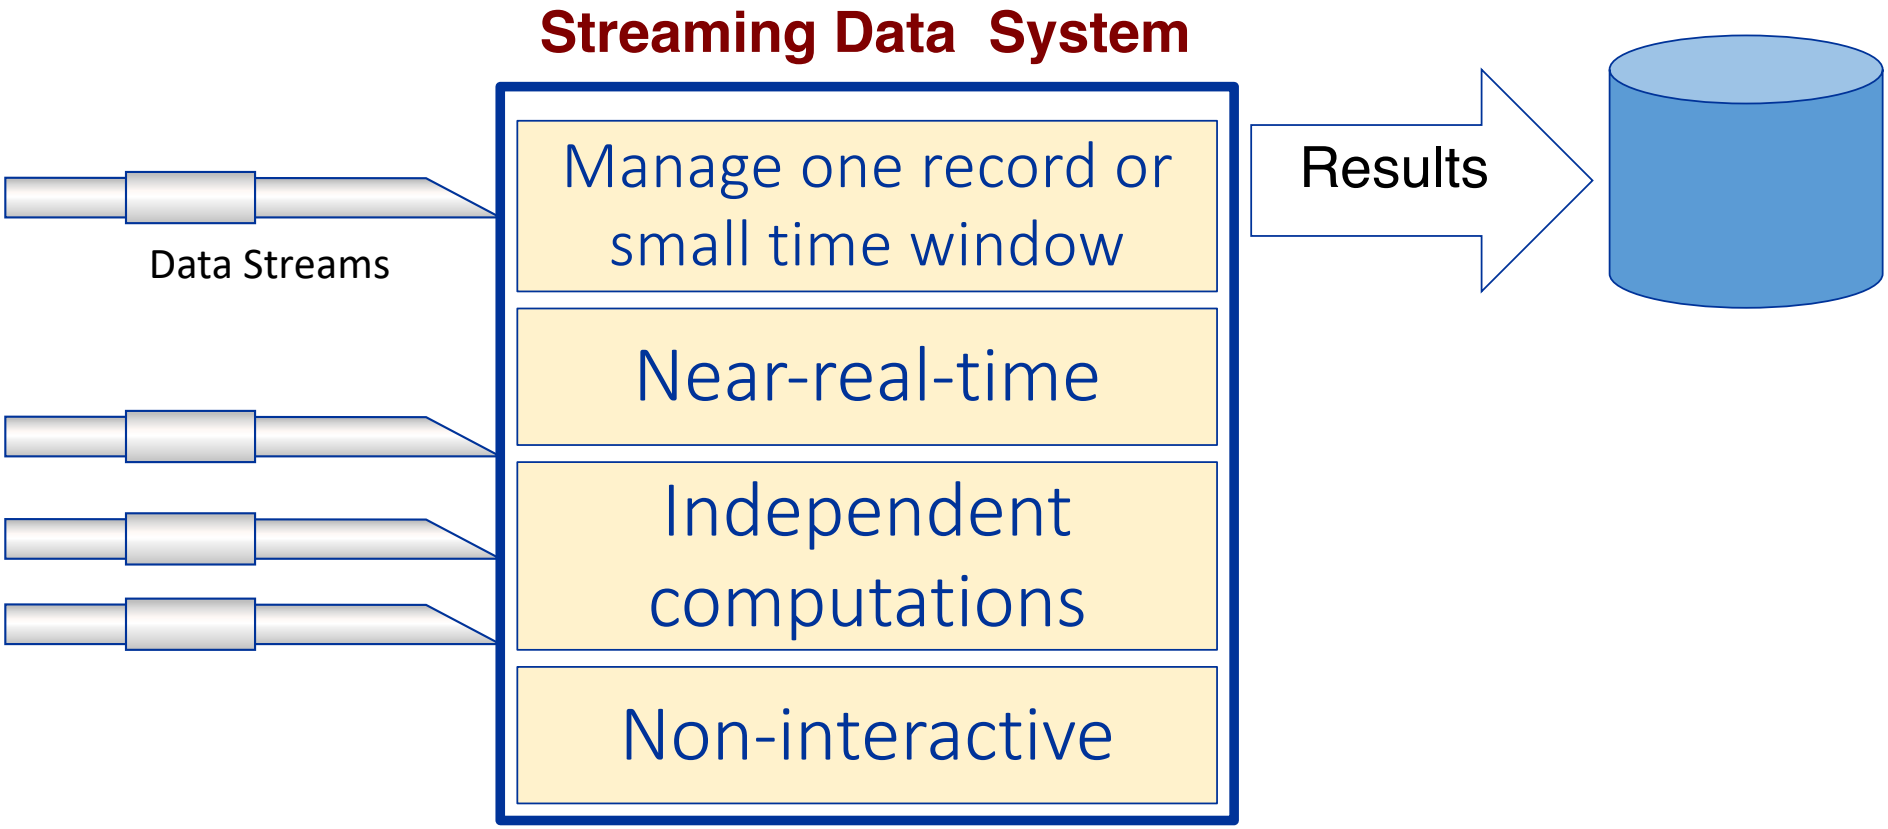

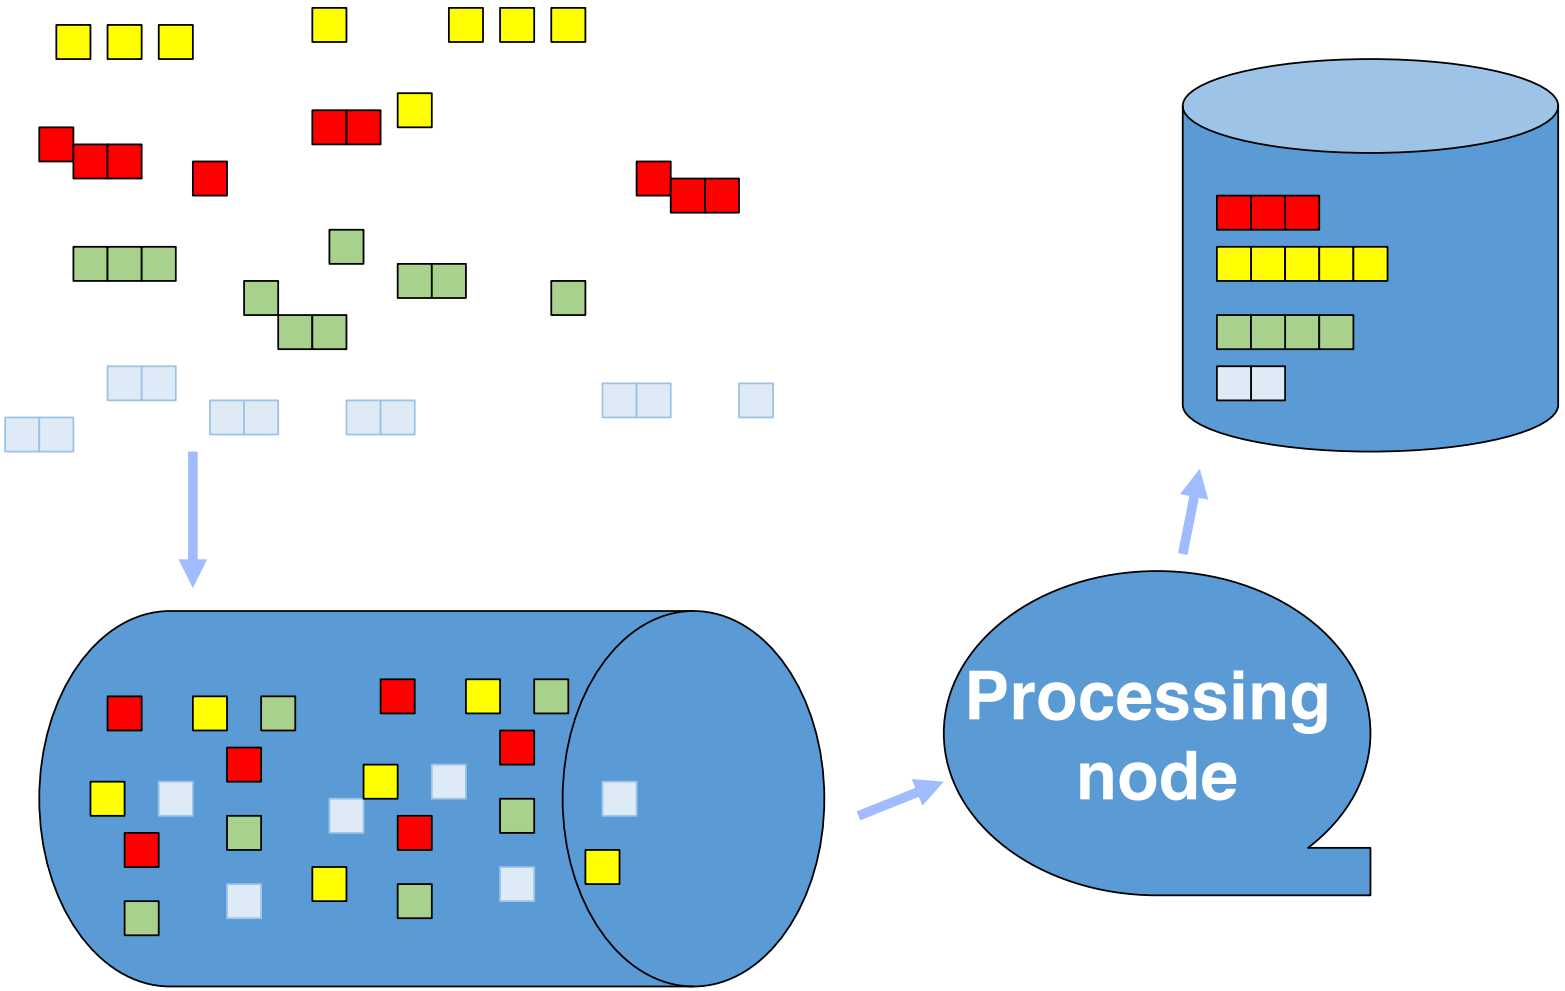
Processing
node

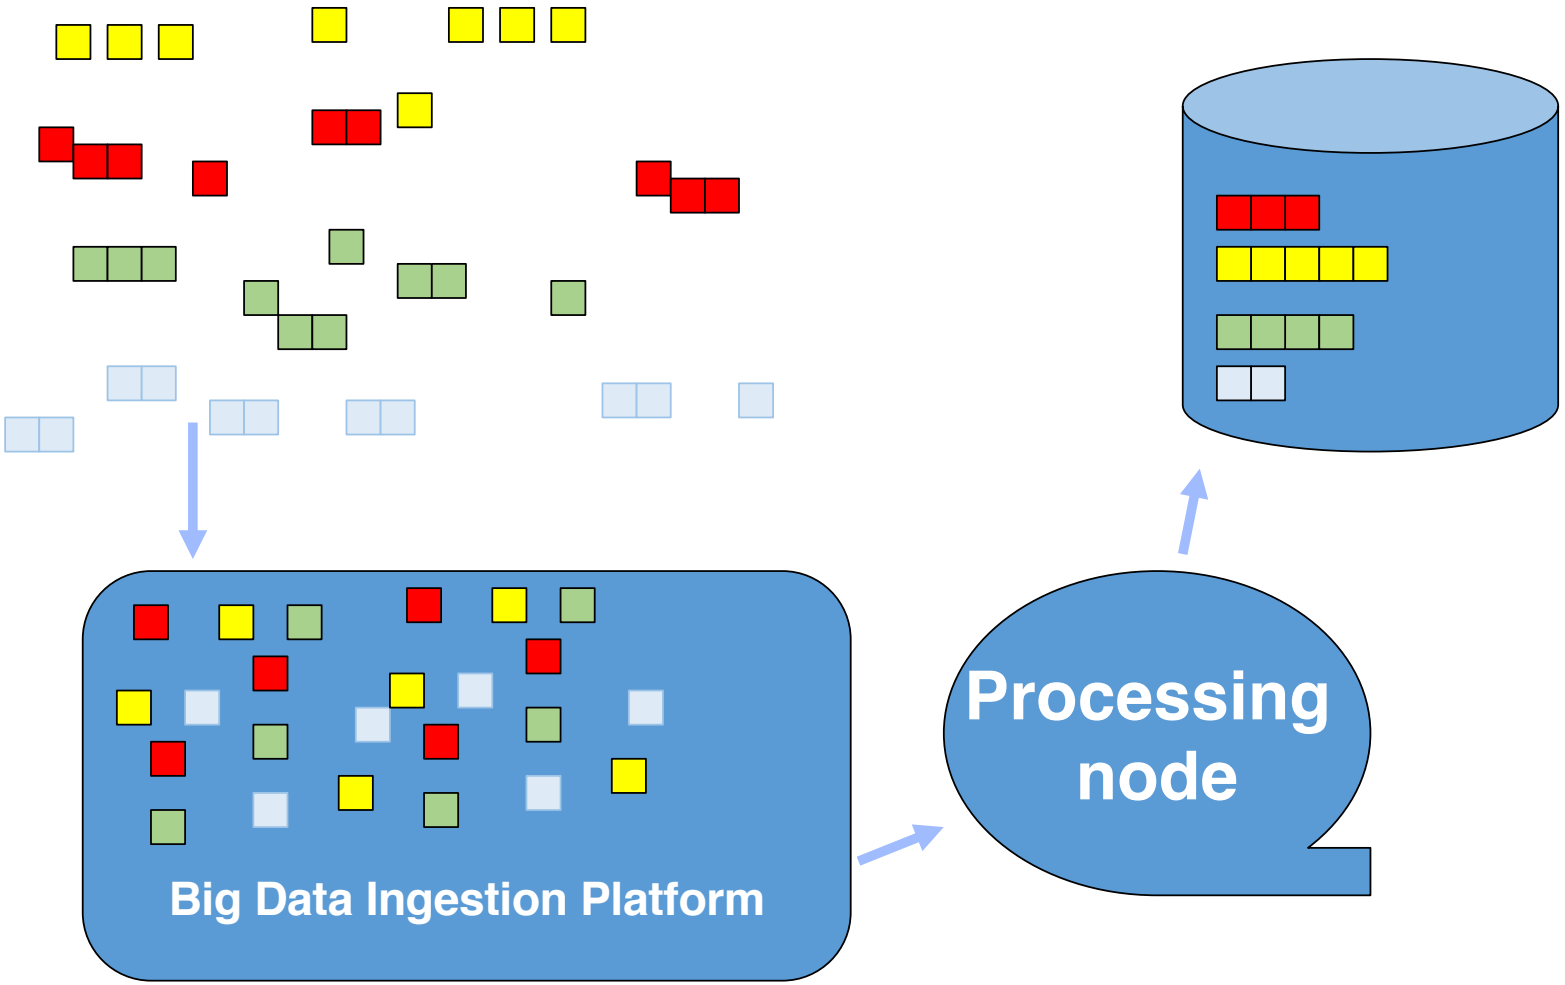
Processing node
Big Data Ingestion Platform

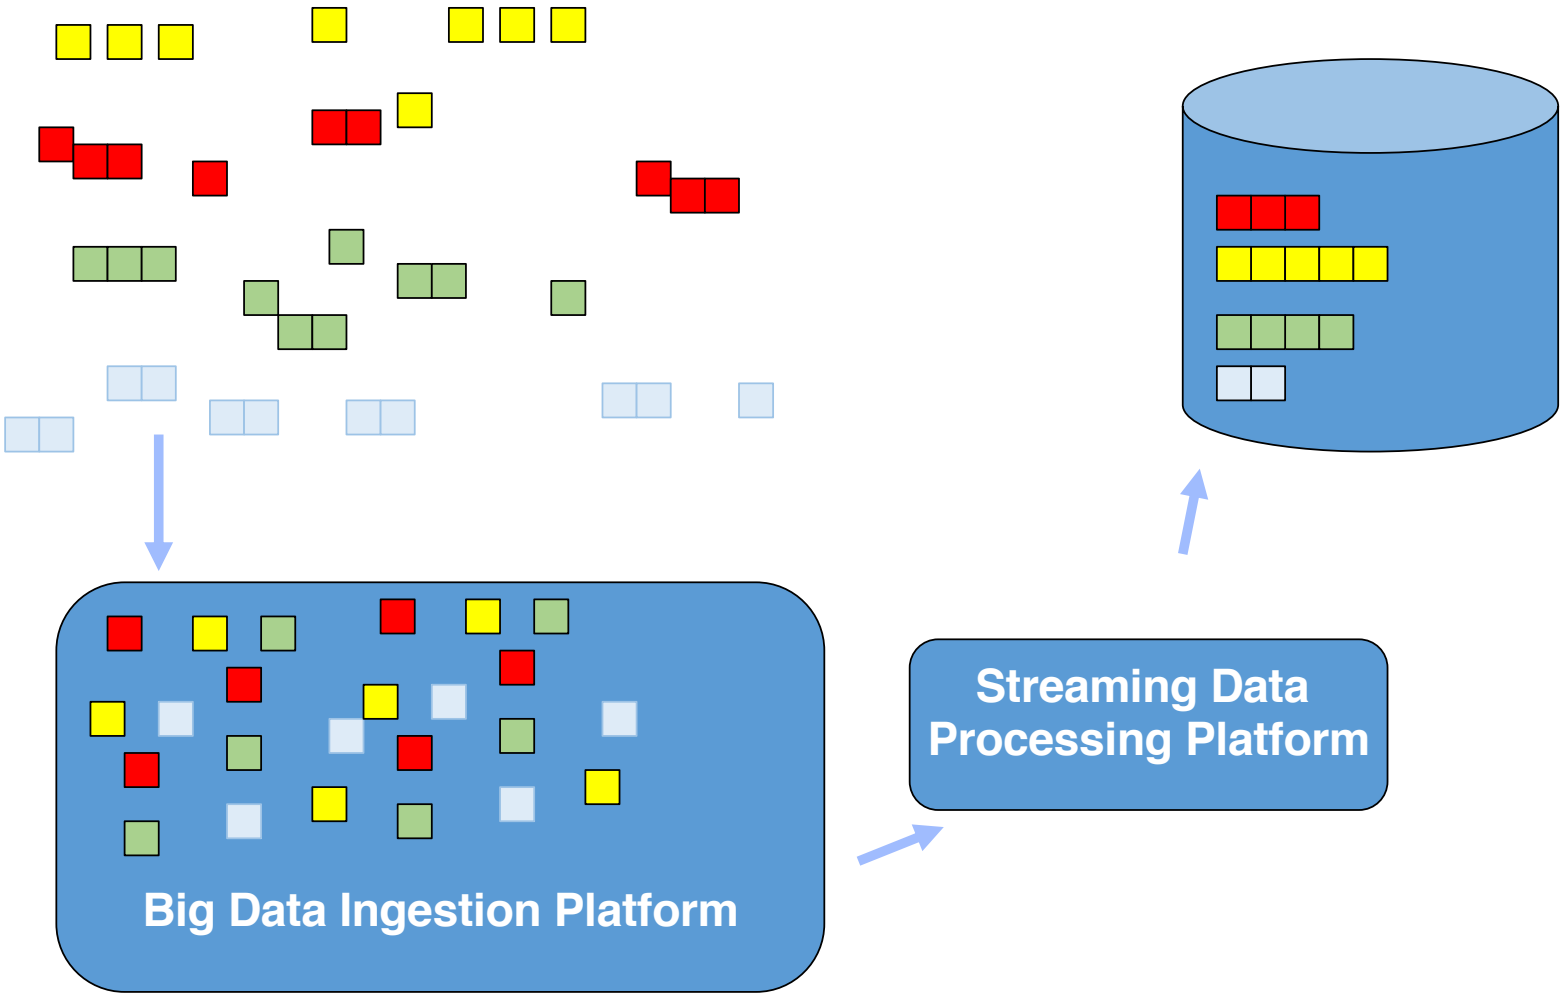
Big Data Ingestion Platform
Streaming Data Processing Platform

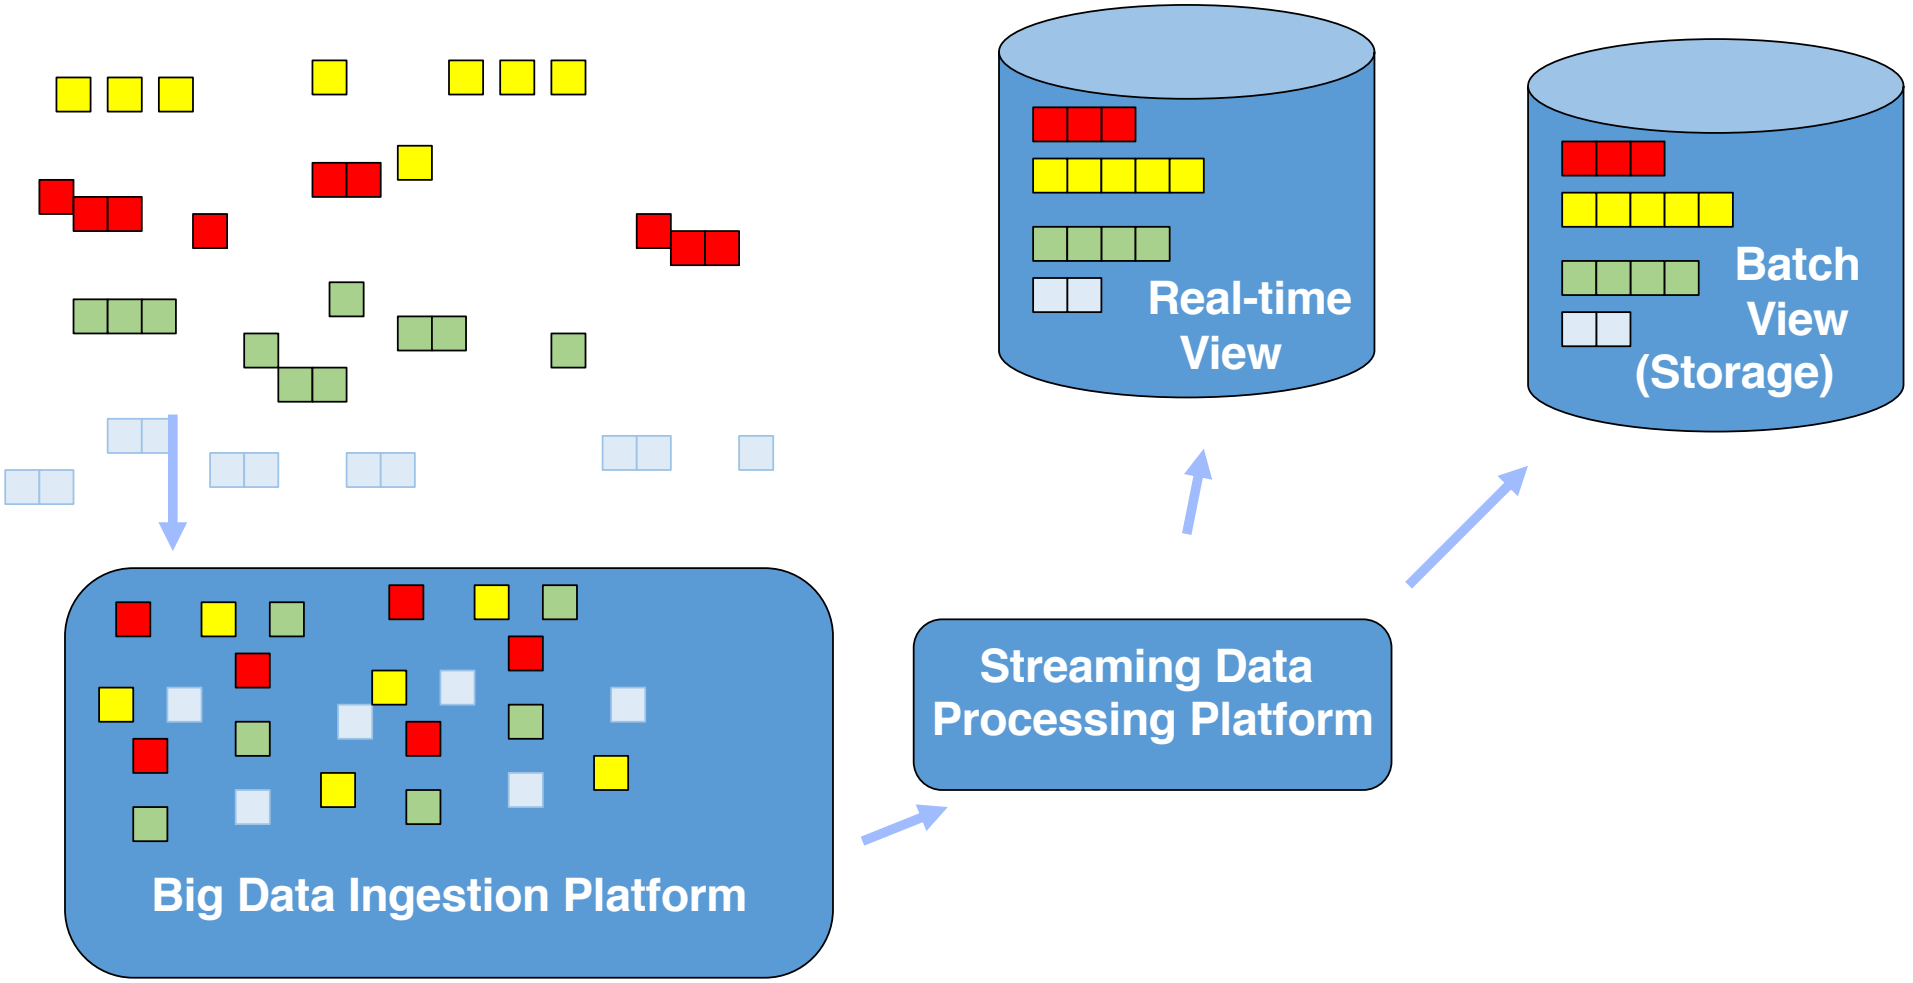
Real-time View
Batch View (Storage)
Streaming Data Processing Platform
Big Data Ingestion Platform

Databases vs. Data Lakes

Short definition of a data lake

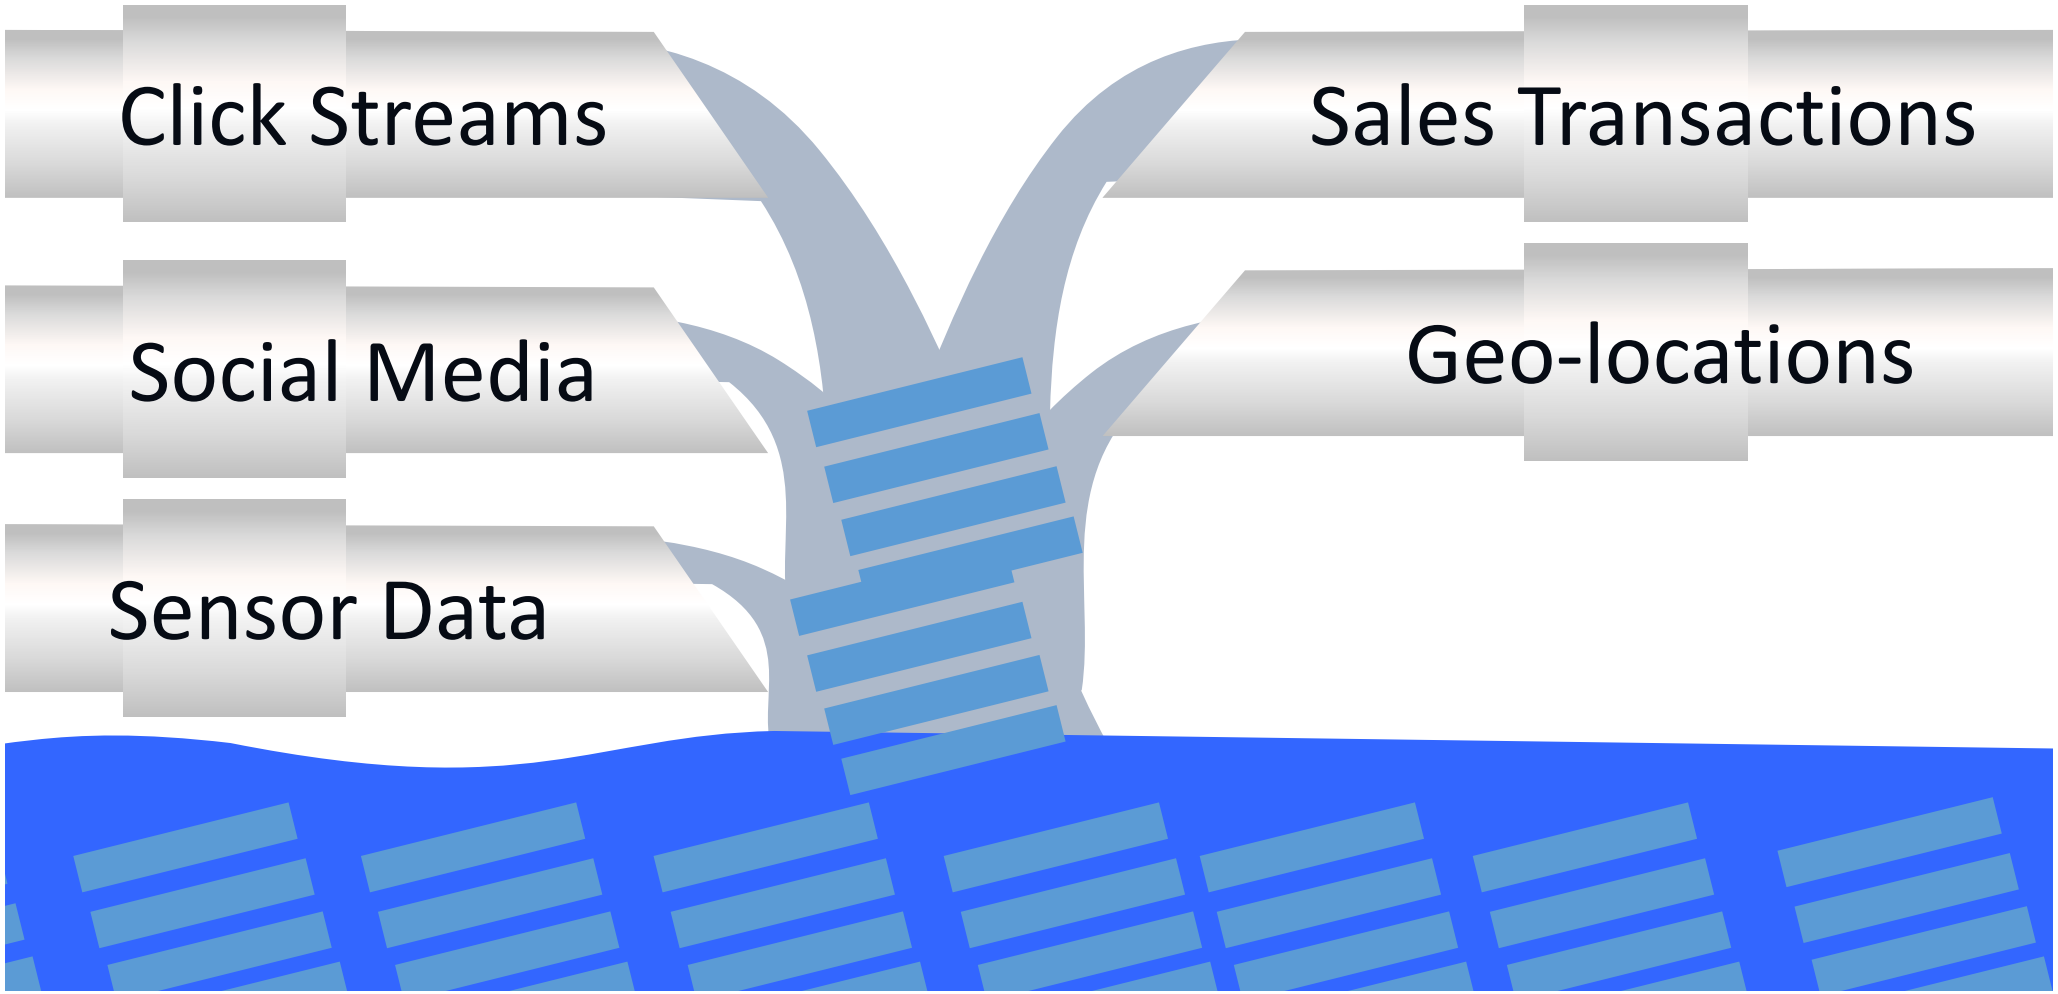

“ If you think of a datamart as a store of bottled water – cleansed and packaged and structured for easy consumption – the data lake is a large body of water in a more natural state. The contents of the data lake stream in from a source to fill the lake, and various users of the lake can come to examine, dive in, or take samples. ”

James Dixon
CTO, Pentaho Corporation

Load data from source

Store raw data

Add data model on read

Data Lake

- Schema-on-read

Data Warehouse

- Schema-on-write
- Transform and structure before load

Hierarchical File System

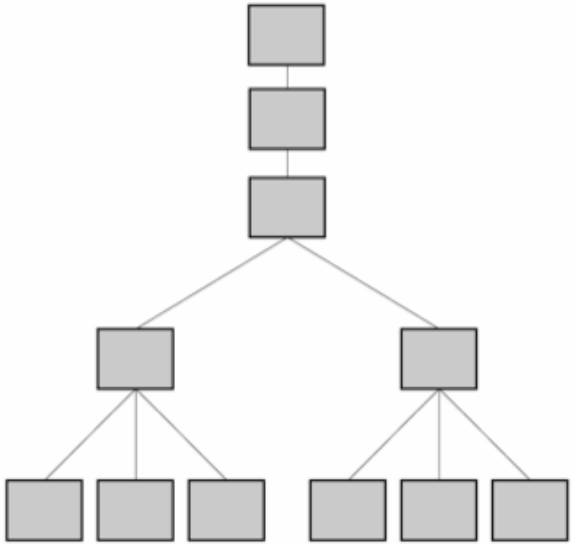

Object storage

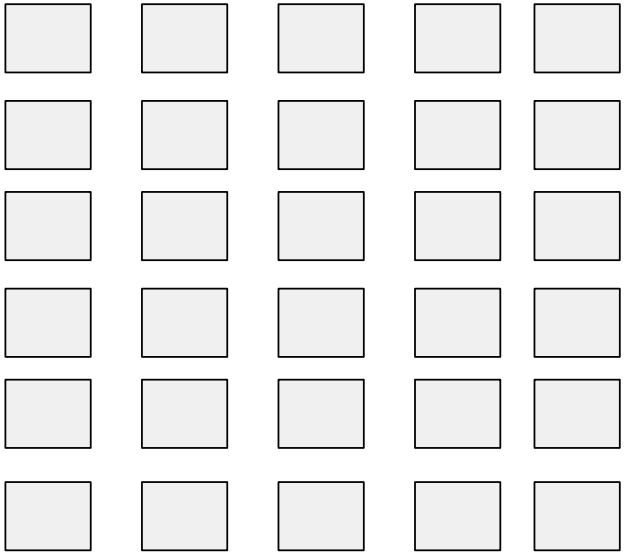

Object Storage

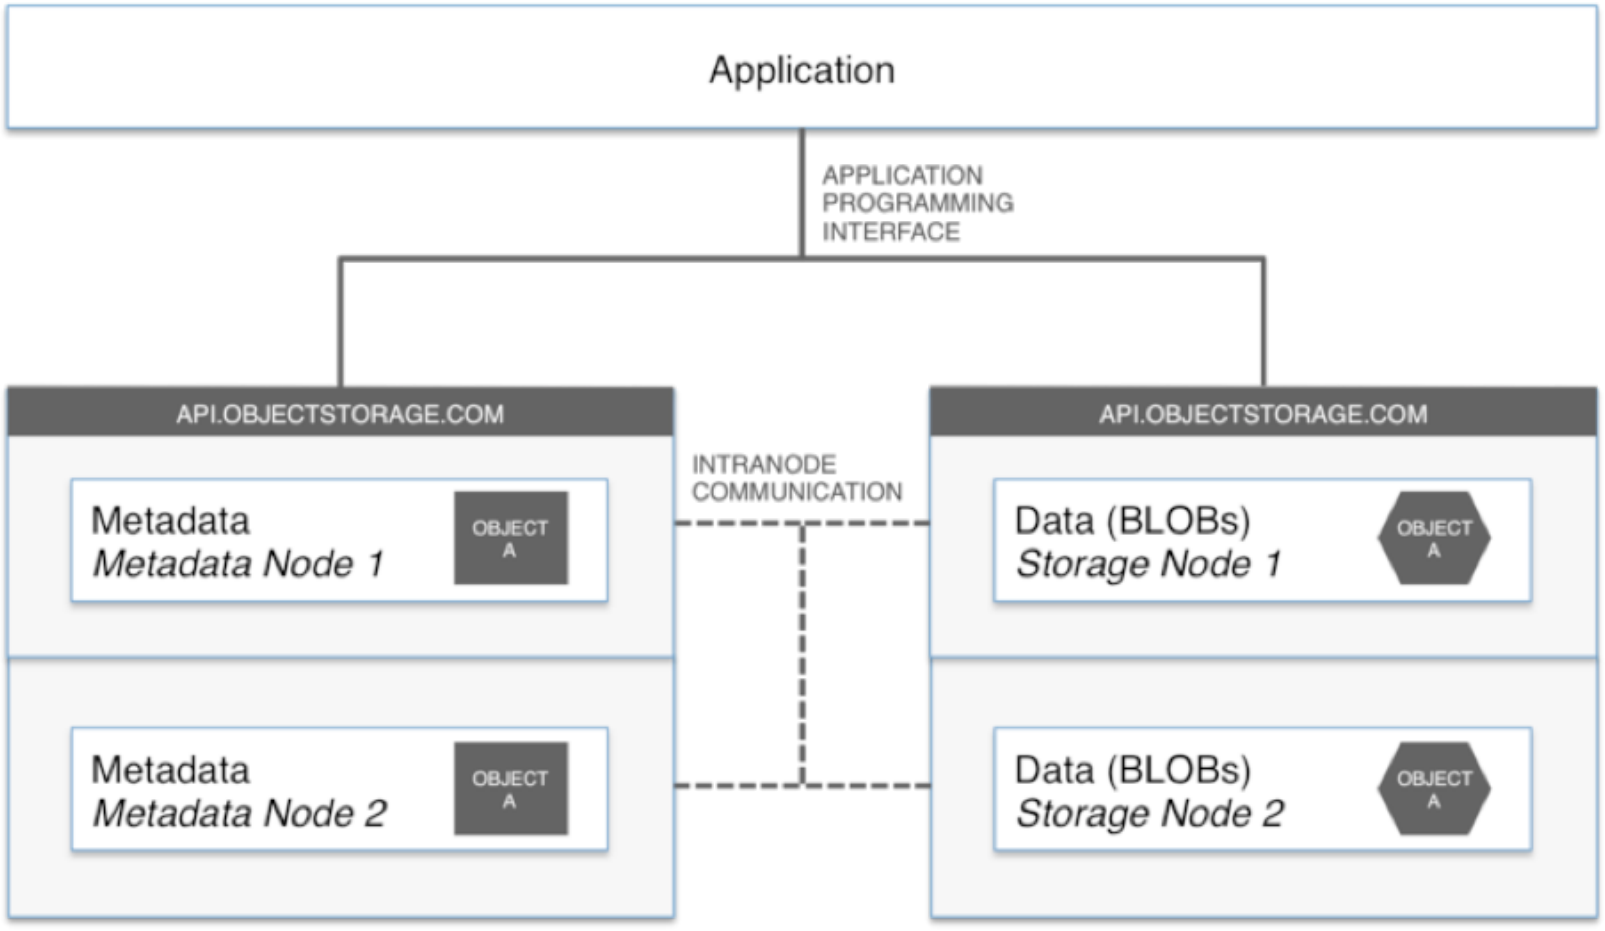

Questions?

ILkay Altintas, Ph.D.
Email: ialtintas@ucsd.edu